

In [1]:

```
pip install pillow
```

Requirement already satisfied: pillow in c:\users\rani_\anaconda6\lib\site-packages (9.4.0)

Note: you may need to restart the kernel to use updated packages.

In [2]:

```
from IPython.display import Image
```

In [3]:

```
Image ("Images/Home builder.jpg" , width=1000, height=500)
```

Out[3]:



KING COUNTY HOUSE SALES ANALYSIS REPORT

Author: Sneha Bhaskar

Date: 24 September 2023

Overview

This project analyzes the King County House Sales Dataset to answer the business case question of what type of housing features to focus when building a house and selling in King County Washington, United States. Utilizing the OSEMN framework for data science, the most valuable feature for a Pro Home Builder Company would be living area square footage and zip code. Zip codes provide contextual information about

property values, local demand, and building regulations, allowing builders to choose locations and tailor their designs to specific demographics strategically. On the other hand, square footage directly influences construction costs, selling price, design considerations, building grade improvements, and the house price rises as well.

Business Problem

The Pro Home Builder Company aims to construct residences in King County Washington, and wants to identify which housing features/areas to focus on for maximum profitability. By understanding which amenities significantly influence home prices, the company can make data-driven choices on the types of houses to build in order to optimize profits

Import standard packages

In [4]:

```
import pandas as pd #data manipulation and analysis. Can create data frames, filter data, import numpy as np #numerical python library for array manipulation, math functions etc. import seaborn as sns #data visualization library based on Matplotlib, with higher-level import matplotlib.pyplot as plt #2D plotting lib. plt.style.use("ggplot") #changes the default style of Matplotlib plots to ggplot.

from statsmodels.formula.api import ols #ordinary Least square (OLS) is a regression mode import statsmodels.api as sm #statsmodel includes various statistical tests, models and f from statsmodels.stats.outliers_influence import variance_inflation_factor #a measure to import scipy.stats as stats #statistical functions from the SciPy lib, which builds on Nu from sklearn.linear_model import LinearRegression #linear model from sklearn.model_selection import train_test_split #function for splitting datasets int from sklearn.metrics import mean_squared_error #function for computing the mean squared from sklearn.preprocessing import OneHotEncoder #preprocessing technique to convert categ from sklearn.model_selection import cross_val_score #function for evaluating a model usin from sklearn.model_selection import KFold #K-Fold cross validator from itertools import combinations #itertools is used for generating all possible combina
```

Kings County House Dataset

In [5]:

```
df = pd.read_csv ('Data/kc_house_data.csv')
df
```

Out[5]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	w
0	7129300520	10/13/2014	221900	3	1.00	1180	5650	1.0	
1	6414100192	12/09/2014	538000	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000	2	1.00	770	10000	1.0	
3	2487200875	12/09/2014	604000	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000	3	2.00	1680	8080	1.0	
...
21592	263000018	5/21/2014	360000	3	2.50	1530	1131	3.0	
21593	6600060120	2/23/2015	400000	4	2.50	2310	5813	2.0	
21594	1523300141	6/23/2014	402101	2	0.75	1020	1350	2.0	
21595	291310100	1/16/2015	400000	3	2.50	1600	2388	2.0	
21596	1523300157	10/15/2014	325000	2	0.75	1020	1076	2.0	

21597 rows × 21 columns

In [6]:

```
#check for duplicates
df.duplicated(keep = False).sum()
```

Out[6]:

0

In [7]:

```
#date range for this dataset
df['date'] = pd.to_datetime(df['date'], format='%m/%d/%Y')
is_sorted = df['date'].is_monotonic_increasing
date_range = pd.date_range(start=df['date'].min(), end=df['date'].max())
print(date_range)
```

```
DatetimeIndex(['2014-05-02', '2014-05-03', '2014-05-04', '2014-05-05',
                 '2014-05-06', '2014-05-07', '2014-05-08', '2014-05-09',
                 '2014-05-10', '2014-05-11',
                 ...
                 '2015-05-18', '2015-05-19', '2015-05-20', '2015-05-21',
                 '2015-05-22', '2015-05-23', '2015-05-24', '2015-05-25',
                 '2015-05-26', '2015-05-27'],
                dtype='datetime64[ns]', length=391, freq='D')
```

OBTAİN

This step involves understanding stakeholder requirements, gathering information on the problem, and finally sourcing data we think will be necessary for solving this problem

In [8]:

```
#Removing columns that is not required from this dataframe for our business problem
df.drop(['id'], axis= 1, inplace = True)
df.head()
```

Out[8]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condit
0	2014-10-13	221900	3	1.00	1180	5650	1.0	NaN	0.0	
1	2014-12-09	538000	3	2.25	2570	7242	2.0	0.0	0.0	
2	2015-02-25	180000	2	1.00	770	10000	1.0	0.0	0.0	
3	2014-12-09	604000	4	3.00	1960	5000	1.0	0.0	0.0	
4	2015-02-18	510000	3	2.00	1680	8080	1.0	0.0	0.0	

In [9]:

```
print("The number of rows is",df.shape[0])
print('The number of columns is',df.shape[1])
```

The number of rows is 21597

The number of columns is 20

SCRUB

This step focuses on cleaning the data, which involves handling missing values, removing outliers, and converting data types and many more. Data cleaning is crucial because the quality of data affects the quality of the final model produced

In [10]:

```
df.info()
#all of the columns are in integer/float format except "date" and "sqft_basement" are in

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 20 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   date              21597 non-null   datetime64[ns]
 1   price              21597 non-null   int64  
 2   bedrooms           21597 non-null   int64  
 3   bathrooms          21597 non-null   float64 
 4   sqft_living        21597 non-null   int64  
 5   sqft_lot            21597 non-null   int64  
 6   floors              21597 non-null   float64 
 7   waterfront          19221 non-null   float64 
 8   view                21534 non-null   float64 
 9   condition           21597 non-null   int64  
 10  grade               21597 non-null   int64  
 11  sqft_above          21597 non-null   int64  
 12  sqft_basement       21597 non-null   object  
 13  yr_built            21597 non-null   int64  
 14  yr_renovated        17755 non-null   float64 
 15  zipcode             21597 non-null   int64  
 16  lat                 21597 non-null   float64 
 17  long                21597 non-null   float64 
 18  sqft_living15       21597 non-null   int64  
 19  sqft_lot15          21597 non-null   int64  
dtypes: datetime64[ns](1), float64(7), int64(11), object(1)
memory usage: 3.3+ MB
```

In [11]:

```
df.describe()
#Bedrooms has an outliers of 33 rooms
#waterfront, condition, grade Looks like a categorical data
```

Out[11]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
count	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04	21597.000000
mean	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04	1.494096
std	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04	0.539683
min	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000
25%	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	1.000000
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04	2.000000
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000

In [12]:

```
#check for null values
df.isna().sum()
#waterfront and yr_renovated are the only columns with null values
```

Out[12]:

date	0
price	0
bedrooms	0
bathrooms	0
sqft_living	0
sqft_lot	0
floors	0
waterfront	2376
view	63
condition	0
grade	0
sqft_above	0
sqft_basement	0
yr_built	0
yr_renovated	3842
zipcode	0
lat	0
long	0
sqft_living15	0
sqft_lot15	0
dtype:	int64

FILL null values for waterfront

In [13]:

```
print(df.waterfront.isnull().sum())
print(df['waterfront'].value_counts())
```

2376	
0.0	19075
1.0	146
Name: waterfront, dtype: int64	

In [14]:

```
print('Mean = ', df['waterfront'].mean())
print('Median = ', df['waterfront'].median())
```

Mean = 0.007595858696217679
Median = 0.0

In [15]:

```
# Replace all the missing values in the waterfront column with the Median
df['waterfront'].fillna(df['waterfront'].median(), inplace=True)
df['waterfront'].isnull().sum()
```

Out[15]:

0

In [16]:

df

Out[16]:

living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_built
1180	5650	1.0	0.0	0.0	3	7	1180	0	1955
2570	7242	2.0	0.0	0.0	3	7	2170	400	1951
770	10000	1.0	0.0	0.0	3	6	770	0	1933
1960	5000	1.0	0.0	0.0	5	7	1050	910	1965
1680	8080	1.0	0.0	0.0	3	8	1680	0	1987
...
1530	1131	3.0	0.0	0.0	3	8	1530	0	2009
2310	5813	2.0	0.0	0.0	3	8	2310	0	2014
1020	1350	2.0	0.0	0.0	3	7	1020	0	2009
1600	2388	2.0	0.0	0.0	3	8	1600	0	2004
1020	1076	2.0	0.0	0.0	3	7	1020	0	2008

◀
▶

Create a new column to replace the yr_renovated where 1 = Renovated and 0 = Not renovated

In [17]:

df.yr_renovated.mean()

Out[17]:

83.6367783722895

In [18]:

df.yr_renovated.median()

Out[18]:

0.0

In [19]:

```
df.yr_renovated.fillna(df.yr_renovated.median(), inplace=True)
df.yr_renovated.isnull().sum()
```

Out[19]:

0

In [20]:

```
df['renovated'] = np.where(df['yr_renovated'] == 0, 0, 1)
```

In [21]:

```
#drop the "yr_renovated" column
df.drop('yr_renovated', axis=1, inplace = True)
```

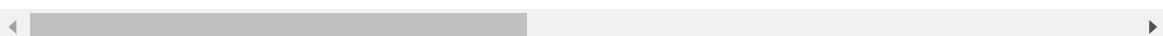
In [22]:

df

Out[22]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
0	2014-10-13	221900	3	1.00	1180	5650	1.0	0.0	0.0	0.0
1	2014-12-09	538000	3	2.25	2570	7242	2.0	0.0	0.0	0.0
2	2015-02-25	180000	2	1.00	770	10000	1.0	0.0	0.0	0.0
3	2014-12-09	604000	4	3.00	1960	5000	1.0	0.0	0.0	0.0
4	2015-02-18	510000	3	2.00	1680	8080	1.0	0.0	0.0	0.0
...
21592	2014-05-21	360000	3	2.50	1530	1131	3.0	0.0	0.0	0.0
21593	2015-02-23	400000	4	2.50	2310	5813	2.0	0.0	0.0	0.0
21594	2014-06-23	402101	2	0.75	1020	1350	2.0	0.0	0.0	0.0
21595	2015-01-16	400000	3	2.50	1600	2388	2.0	0.0	0.0	0.0
21596	2014-10-15	325000	2	0.75	1020	1076	2.0	0.0	0.0	0.0

21597 rows × 20 columns



In [23]:

```
df.renovated.value_counts()
```

Out[23]:

```
0    20853  
1     744  
Name: renovated, dtype: int64
```

Clean "Sqft_Basement" column by replacing '?' with 0 and converting the data type from object to float

In [24]:

```
df.sqft_basement.value_counts()
```

Out[24]:

```
0      12826  
?      454  
600    217  
500    209  
700    208  
...  
1920    1  
3480    1  
2730    1  
2720    1  
248     1  
Name: sqft_basement, Length: 304, dtype: int64
```

In [25]:

```
#replace the '?' in the basement column with 0  
df['sqft_basement'] = df['sqft_basement'].replace(to_replace='?', value='0.0')
```

In [26]:

```
df.sqft_basement.value_counts()
```

Out[26]:

```
0      12826  
0.0    454  
600    217  
500    209  
700    208  
...  
1920    1  
3480    1  
2730    1  
2720    1  
248     1  
Name: sqft_basement, Length: 304, dtype: int64
```

In [27]:

```
#convert the data type of 'sqft_basement' from object to float
df['sqft_basement'] = df['sqft_basement'].astype(float)
```

In [28]:

```
#confirm we don't have any null values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   date             21597 non-null   datetime64[ns]
 1   price            21597 non-null   int64  
 2   bedrooms          21597 non-null   int64  
 3   bathrooms         21597 non-null   float64
 4   sqft_living       21597 non-null   int64  
 5   sqft_lot          21597 non-null   int64  
 6   floors            21597 non-null   float64
 7   waterfront        21597 non-null   float64
 8   view              21534 non-null   float64
 9   condition         21597 non-null   int64  
 10  grade             21597 non-null   int64  
 11  sqft_above        21597 non-null   int64  
 12  sqft_basement     21597 non-null   float64
 13  yr_built          21597 non-null   int64  
 14  zipcode           21597 non-null   int64  
 15  lat               21597 non-null   float64
 16  long              21597 non-null   float64
 17  sqft_living15     21597 non-null   int64  
 18  sqft_lot15        21597 non-null   int64  
 19  renovated          21597 non-null   int32  
dtypes: datetime64[ns](1), float64(7), int32(1), int64(11)
memory usage: 3.2 MB
```

In [29]:

```
#check the data statistics
df.describe()
```

Out[29]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
count	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04	21597.000000
mean	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04	1.494096
std	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04	0.539683
min	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000
25%	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	1.000000
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04	2.000000
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000

PRICE

In [30]:

```
# Check price distribution
df['price'].div(1000000).hist(color = "yellow")
plt.title('Price Distribution');
plt.xlabel('Price - in Millions :)');
plt.ylabel('Number of Houses');

#the majority of houses are below the price of 2 million
```

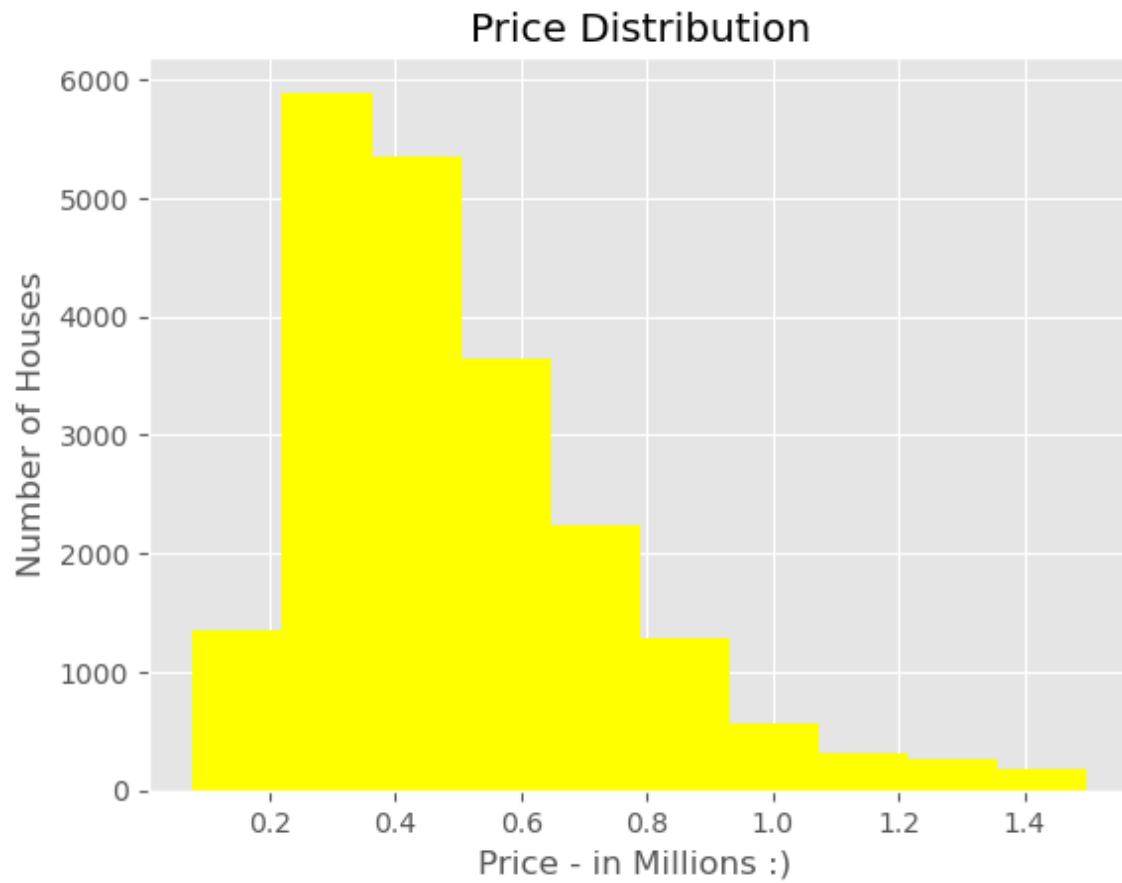


In [31]:

```
#reduce the data to only houses equal to or below 1.5 million
df_og = df.copy()
# remove houses above 1.5 million, boolean mask
df = df[df['price'] <= 1500000]
```

In [32]:

```
# RECHECK price distribution
df['price'].div(1000000).hist(color = "yellow")
plt.title('Price Distribution');
plt.xlabel('Price - in Millions :)');
plt.ylabel('Number of Houses');
```

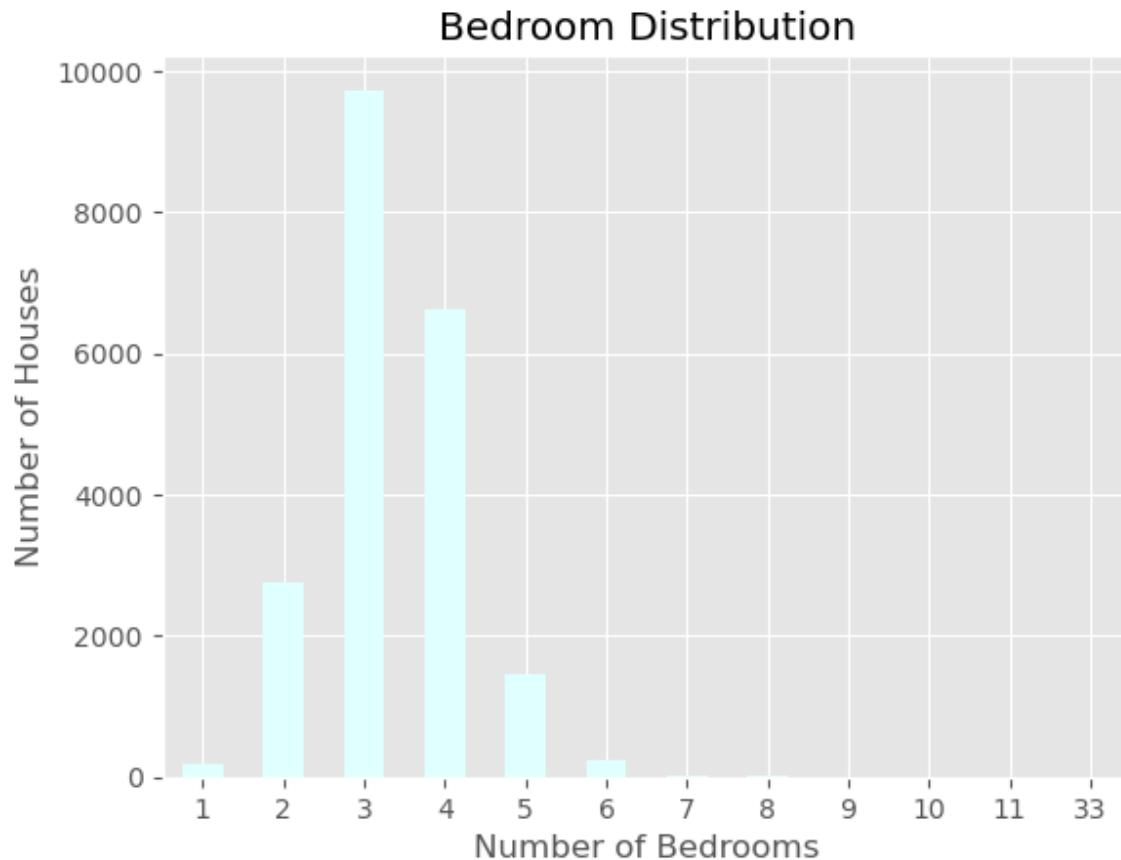


Majority of the houses are in between 0.3 - 0.5 price range

BEDROOMS

In [33]:

```
# Plot bedrooms as categorical - ordinal
df['bedrooms'].value_counts().sort_index().plot(kind = 'bar', color = "lightcyan")
plt.xticks(rotation = 0);
plt.title('Bedroom Distribution');
plt.xlabel('Number of Bedrooms');
plt.ylabel('Number of Houses');
```



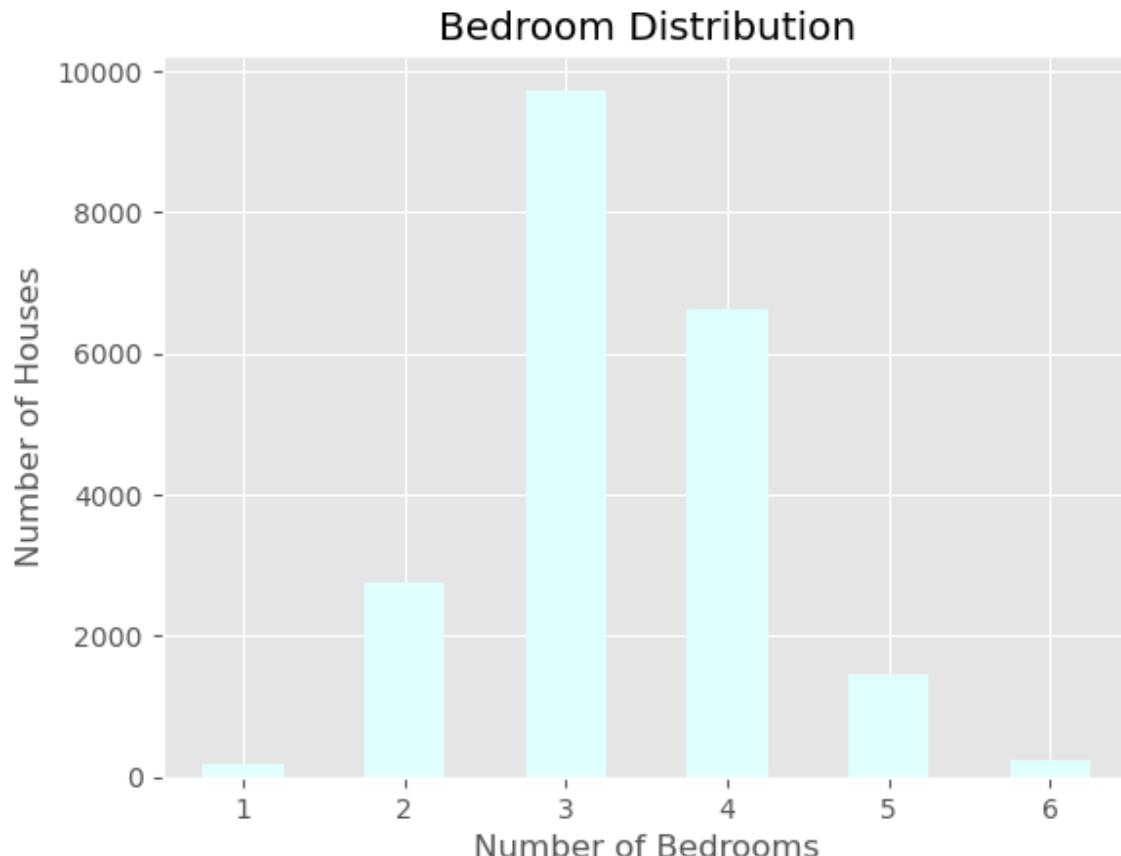
In [34]:

```
#reduce the data to only bedrooms equal to or below 6
df_og = df.copy()
# remove bedrooms above 6, boolean mask
df = df[df['bedrooms'] <=6]
```

In [35]:

```
# Plot bedrooms as categorical - ordinal
df['bedrooms'].value_counts().sort_index().plot(kind = 'bar', color = "lightcyan")
plt.xticks(rotation = 0);
plt.title('Bedroom Distribution');
plt.xlabel('Number of Bedrooms');
plt.ylabel('Number of Houses');

#by removing the 33 outliers for bedrooms , its summarises that houses with 3 or 4 bedro
```

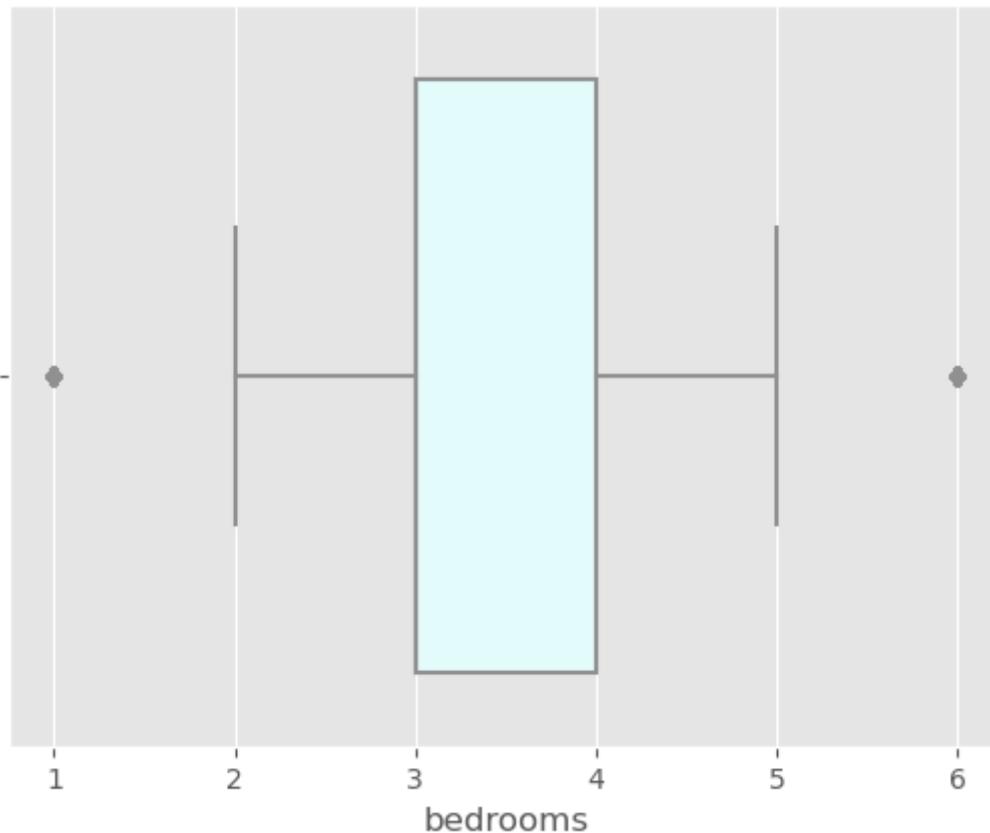


Houses with 3 or 4 bedrooms are the most common in the dataset

In [36]:

```
sns.boxplot(x = df['bedrooms'], color = "lightcyan");
```

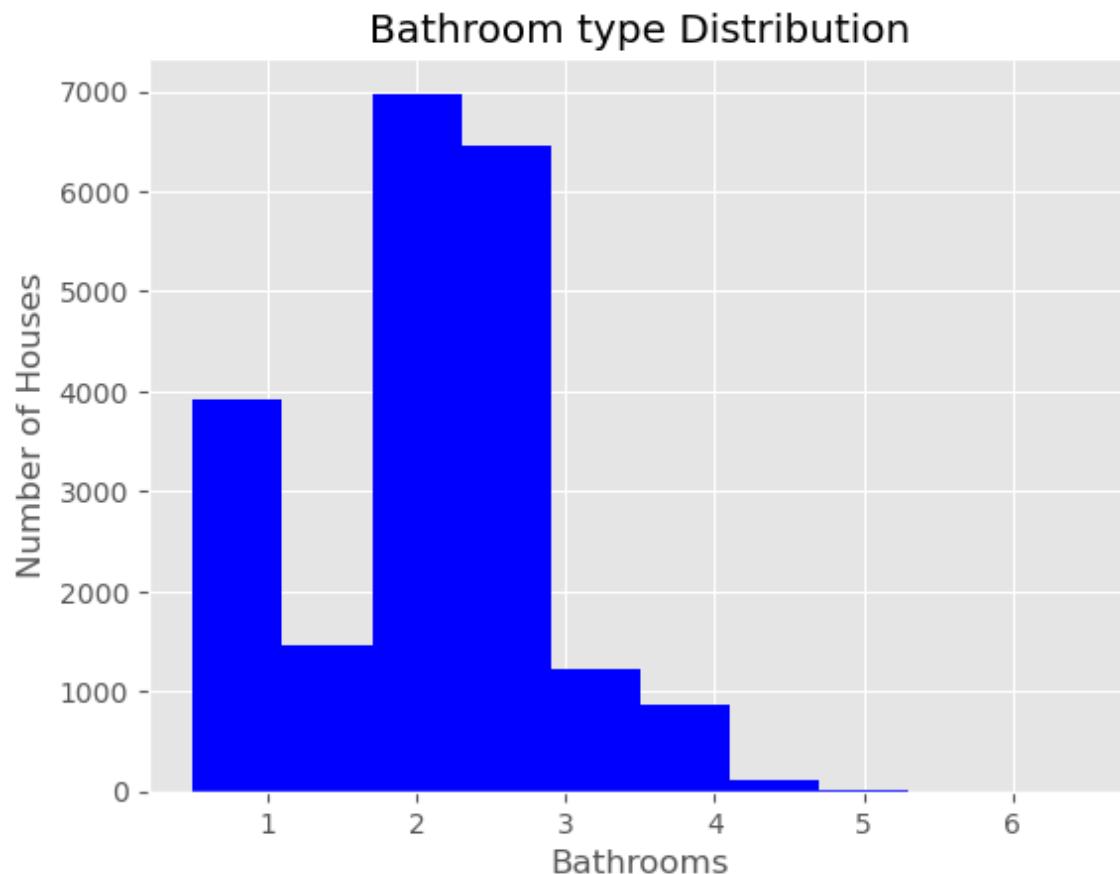
#Boxplot shows are normal distribution



BATHROOMS

In [37]:

```
# Check bathroom types distribution
df['bathrooms'].hist(color = "BLUE")
plt.title('Bathroom type Distribution');
plt.xlabel('Bathrooms');
plt.ylabel('Number of Houses');
```

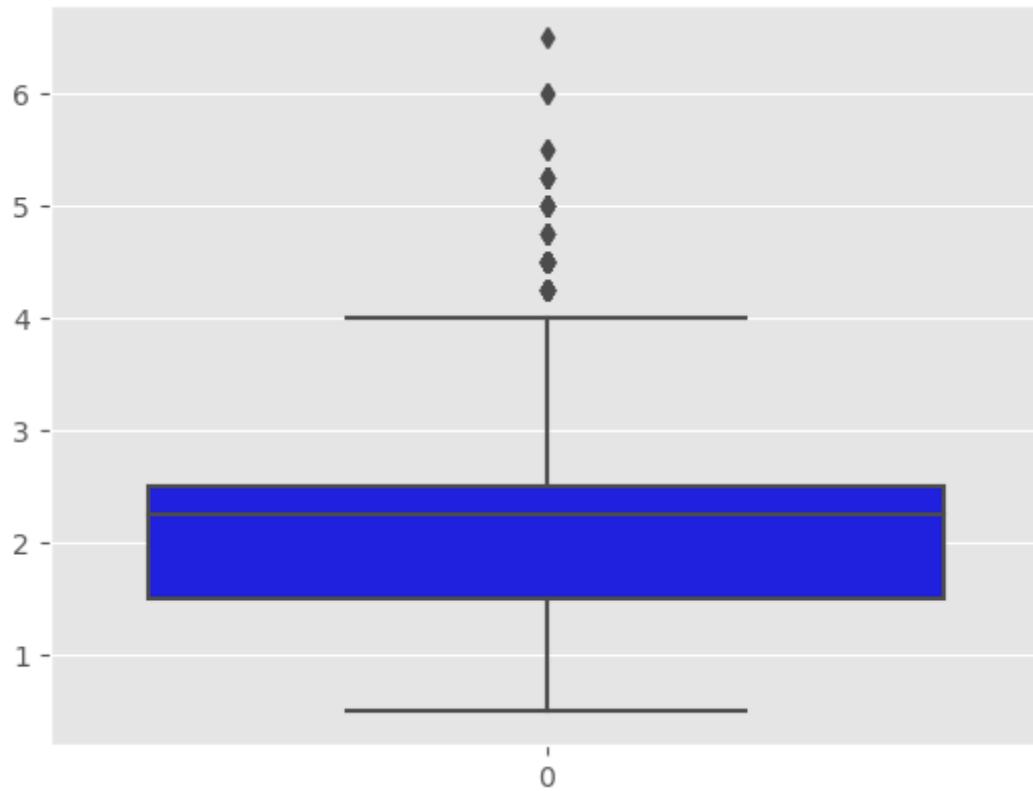


There appears to be an outlier and most of the houses have 2 or 3 bathrooms

In [38]:

```
# Boxplot to detect outliers
sns.boxplot(df['bathrooms'], color = "BLUE");

#there are handful of outliers above 4
```



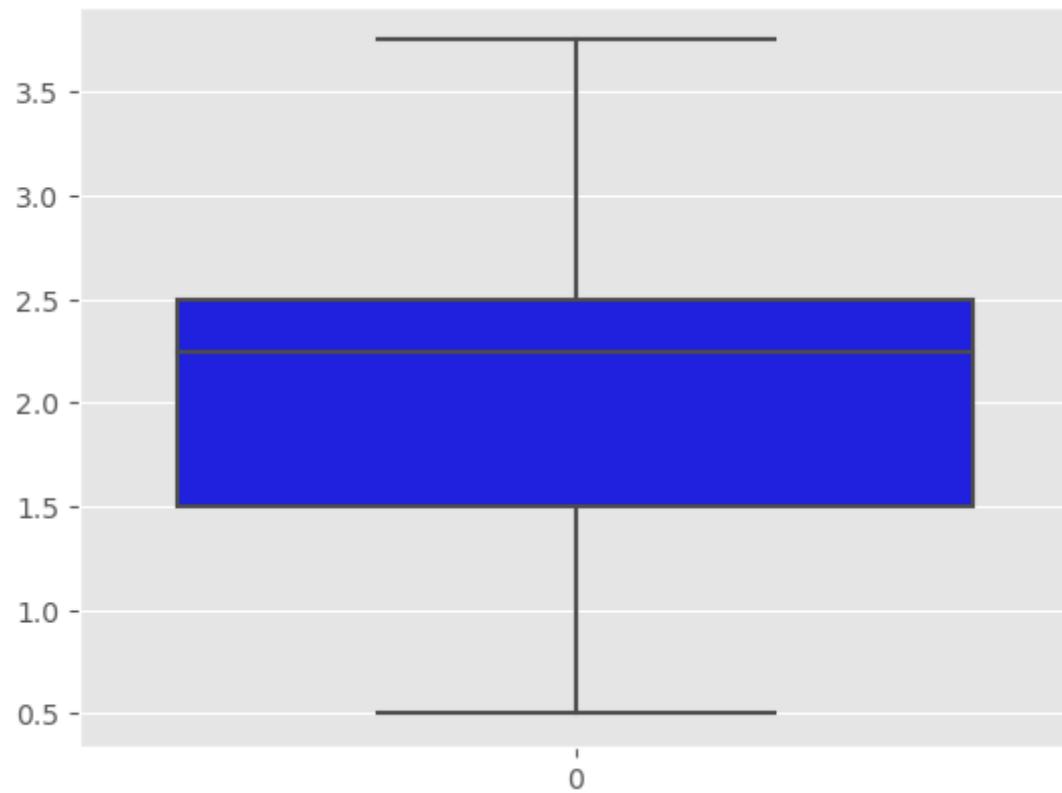
In [39]:

```
# Drop houses that have more than 4 bathrooms
df = df[df['bathrooms'] < 4]
```

In [40]:

```
# Boxplot to detect outliers
sns.boxplot(df['bathrooms'], color = "BLUE");

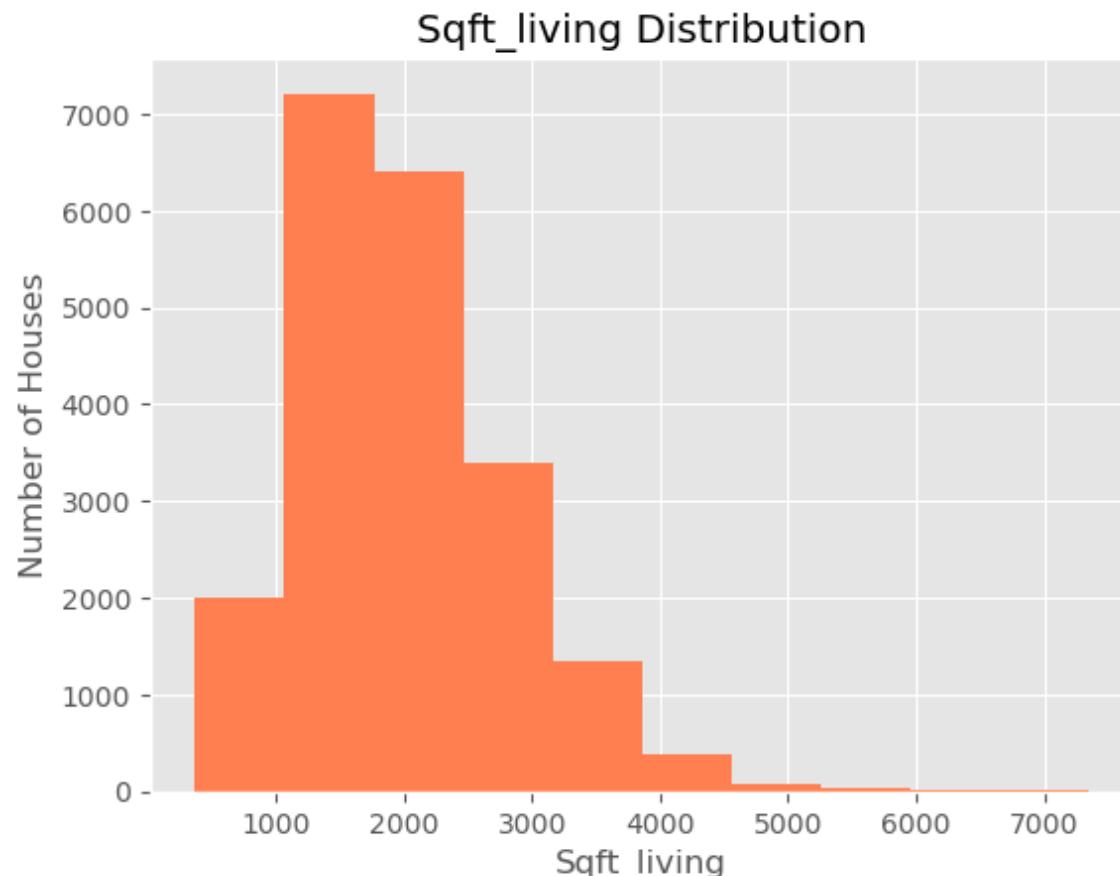
#Boxplot shows are normal distribution
```



SQFT_LIVING

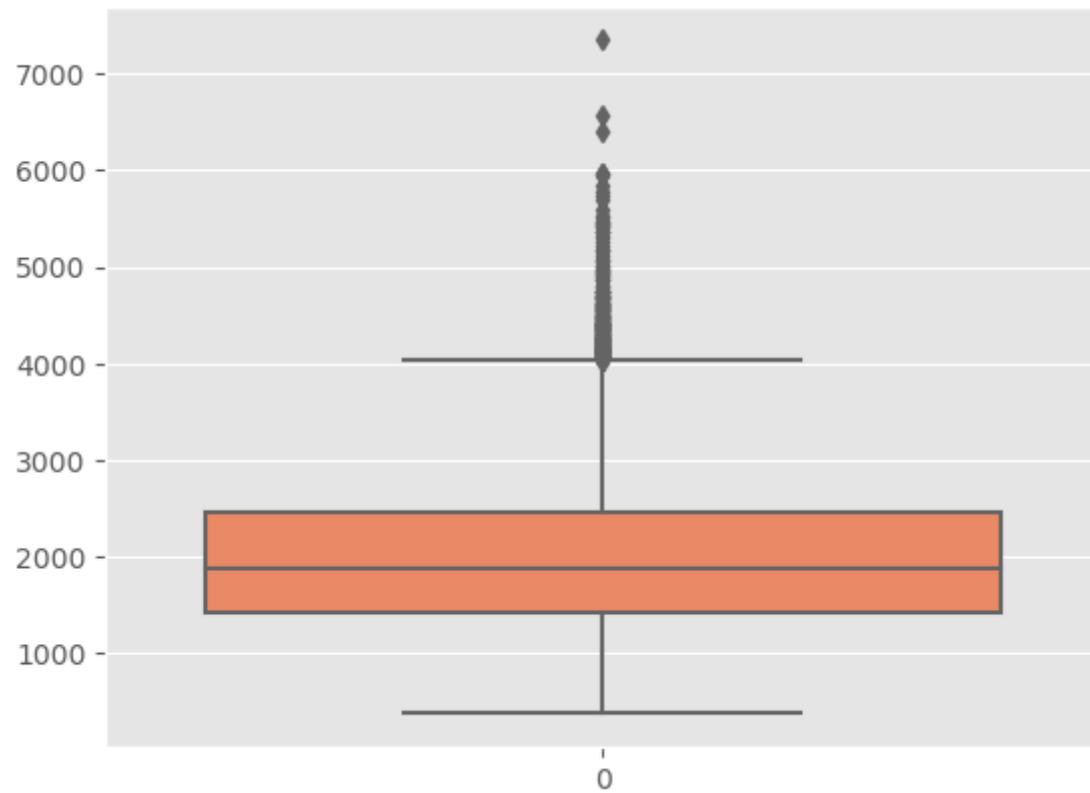
In [41]:

```
# Check sqft_living distribution
df['sqft_living'].hist(color = "coral")
plt.title('Sqft_living Distribution');
plt.xlabel('Sqft_living');
plt.ylabel('Number of Houses');
```



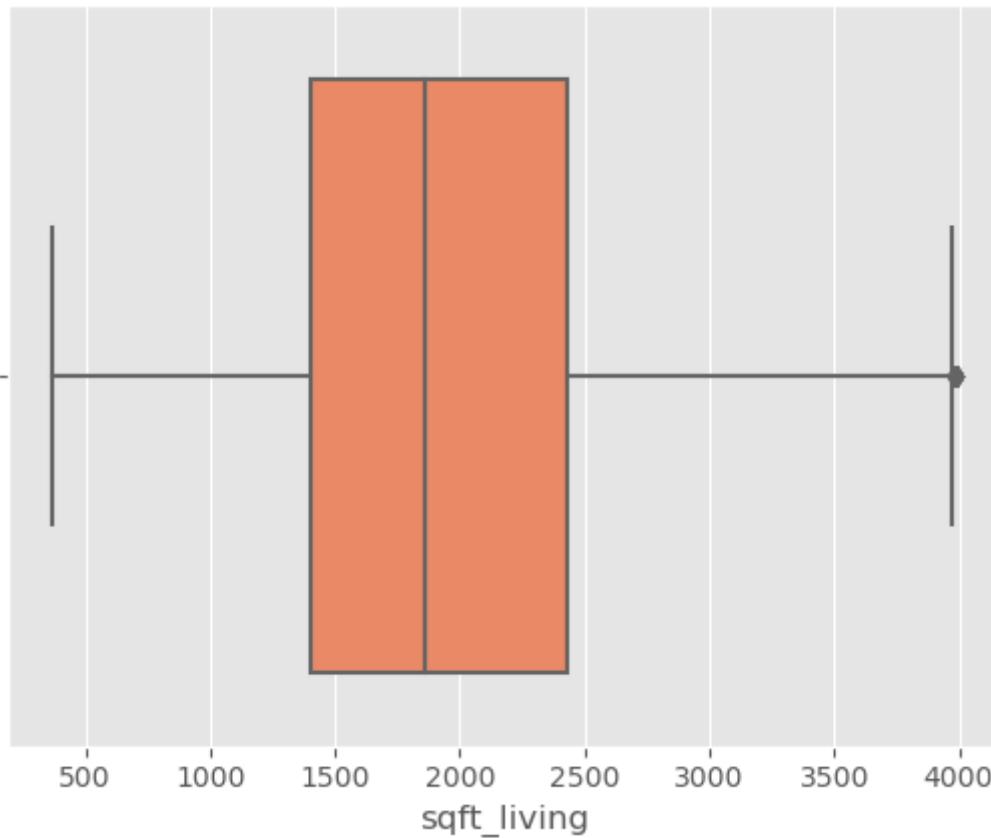
In [42]:

```
# Box plot for outliers  
#There are outliers above 4000 sqft_living.  
sns.boxplot(df['sqft_living'], color = "coral");
```



In [43]:

```
# Drop outliers in sqft_living
df = df[df['sqft_living'] < 4000]
sns.boxplot(x = df['sqft_living'], color = "coral");
```

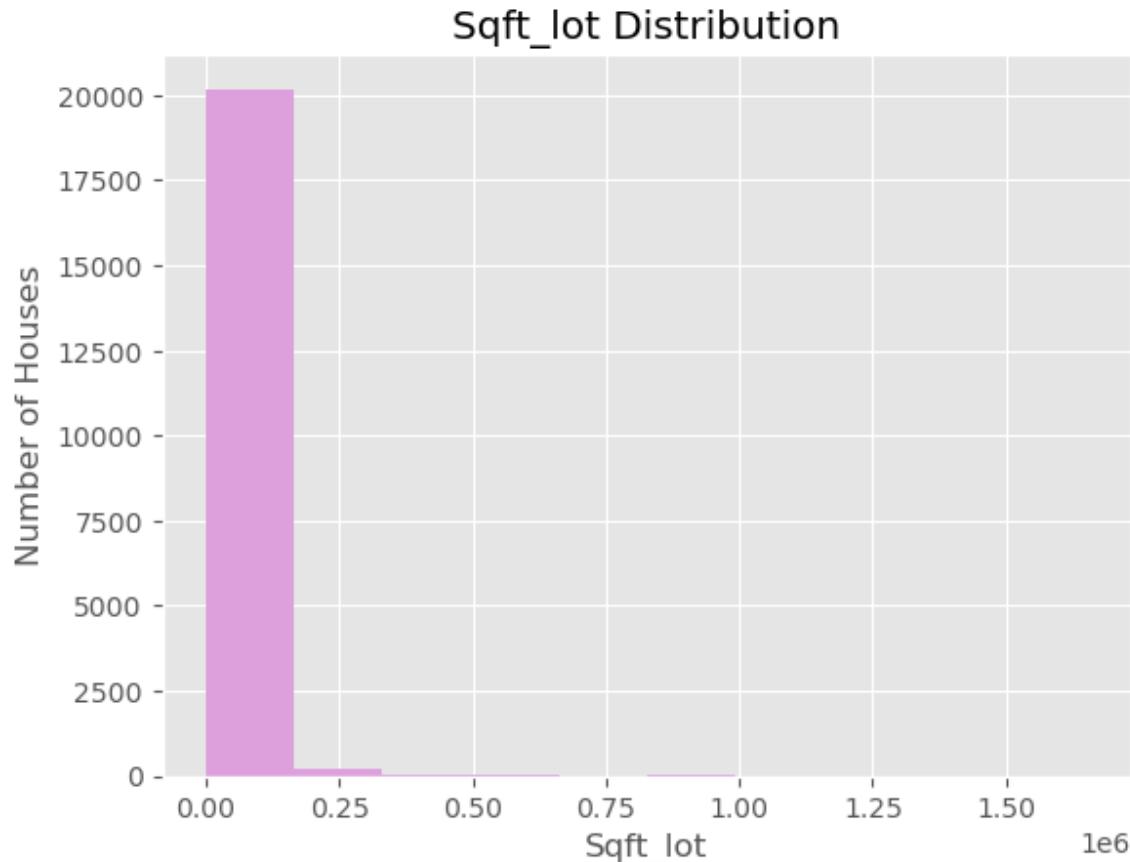


1500 - 2500 sqft_living is most commonly observed footage of the house in this dataset

SQFT_LOT

In [44]:

```
# Check sqft_lot distribution
df['sqft_lot'].hist(color = "plum")
plt.title('Sqft_lot Distribution');
plt.xlabel('Sqft_lot');
plt.ylabel('Number of Houses');
```



In [45]:

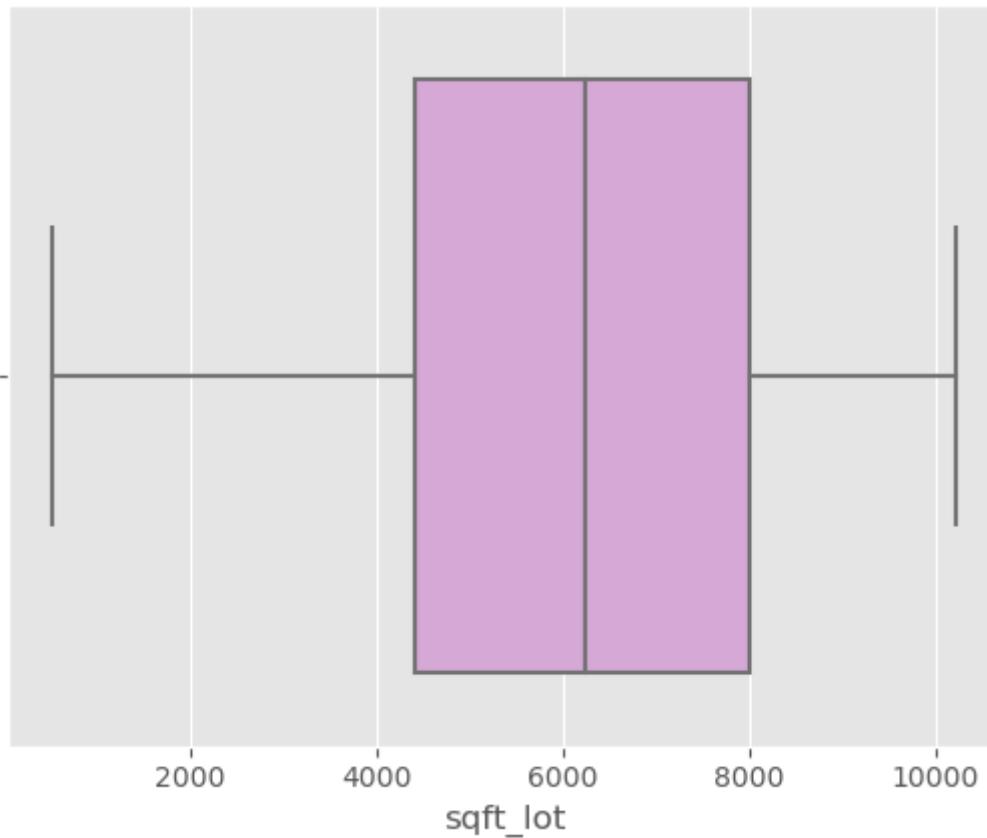
```
df.describe()
```

Out[45]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
count	2.044600e+04	20446.000000	20446.000000	20446.000000	2.044600e+04	20446.000000
mean	4.872377e+05	3.310036	2.033735	1954.156705	1.399462e+04	1.472684
std	2.327866e+05	0.850293	0.676472	724.038899	3.766160e+04	0.536808
min	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000
25%	3.150000e+05	3.000000	1.500000	1400.000000	5.000000e+03	1.000000
50%	4.380000e+05	3.000000	2.000000	1860.000000	7.500000e+03	1.000000
75%	6.050000e+05	4.000000	2.500000	2430.000000	1.023275e+04	2.000000
max	1.500000e+06	6.000000	3.750000	3990.000000	1.651359e+06	3.500000

In [46]:

```
# Drop houses that are above the 75 percentile
df = df[df['sqft_lot'] < 1.023275e+04]
sns.boxplot(x = df['sqft_lot'], color = "plum");
```

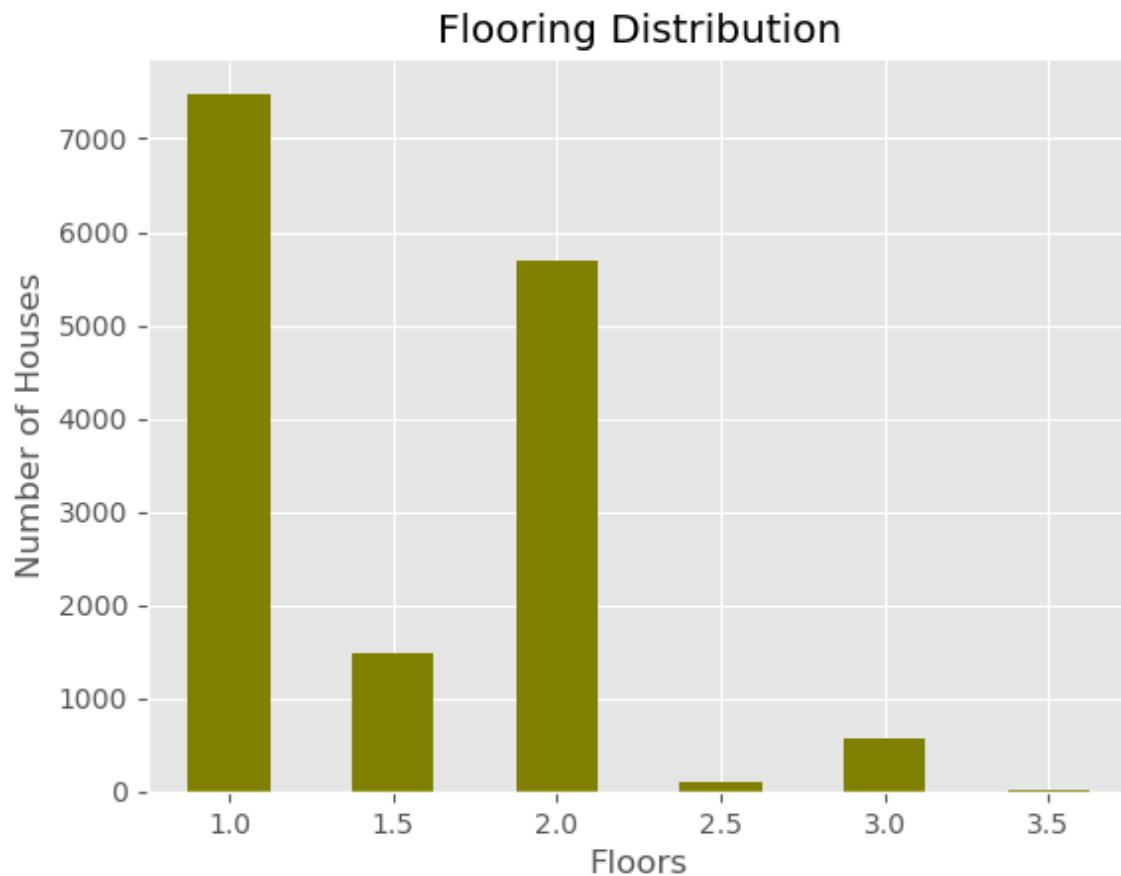


The majority of houses have between 4000 and 8000 sqft_lot.

FLOORS

In [47]:

```
# Plot a bar chart to check distribution of flooring
df['floors'].value_counts().sort_index().plot(kind = 'bar', color = "OLIVE")
plt.xticks(rotation = 0)
plt.title('Flooring Distribution');
plt.xlabel('Floors');
plt.ylabel('Number of Houses');
```

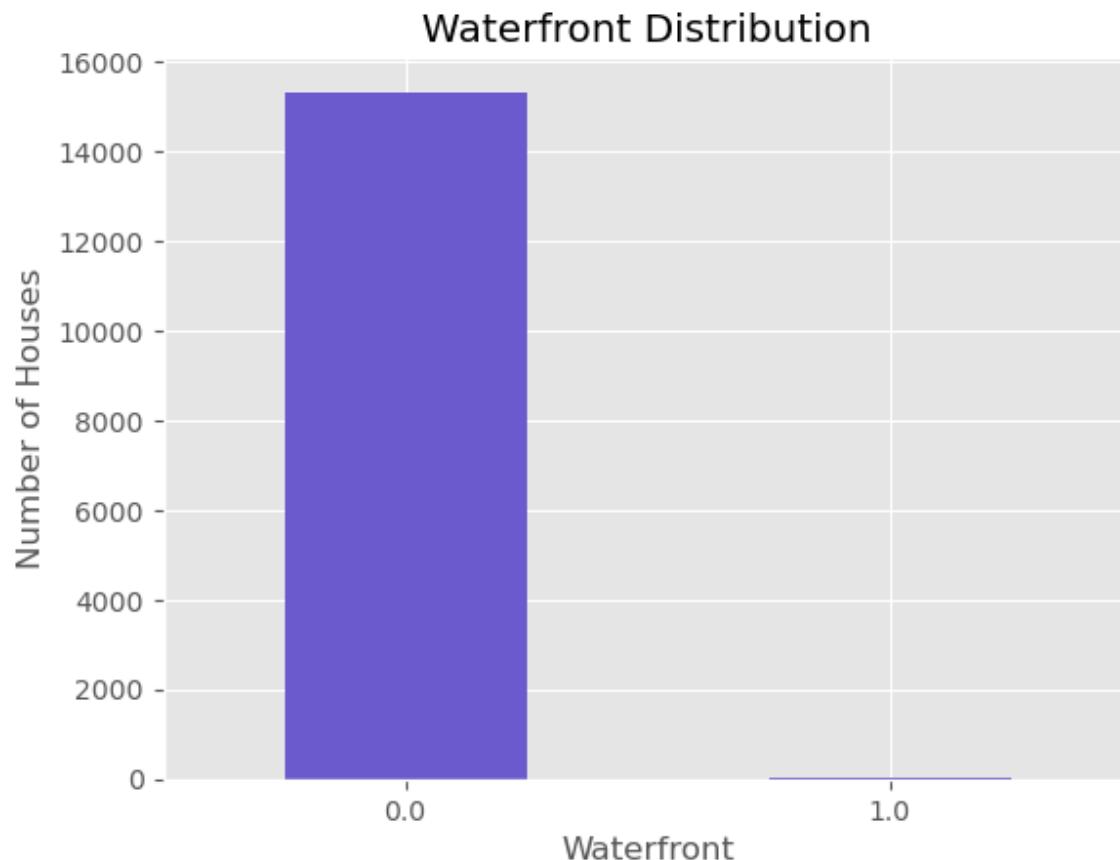


Most houses have either 1 or 2 floors.

WATERFRONT

In [48]:

```
# Plot a bar chart for waterfront
df['waterfront'].value_counts().sort_index().plot(kind = 'bar', color = "slateblue")
plt.xticks(rotation = 0)
plt.title('Waterfront Distribution');
plt.xlabel('Waterfront');
plt.ylabel('Number of Houses');
```



Waterfront is categorical as it has two categories. Also, the majority of the homes are not waterfront. 99% of houses do not have waterfront.

In [49]:

```
#checking data types  
df.dtypes
```

Out[49]:

```
date          datetime64[ns]  
price         int64  
bedrooms      int64  
bathrooms     float64  
sqft_living   int64  
sqft_lot      int64  
floors        float64  
waterfront    float64  
view          float64  
condition     int64  
grade         int64  
sqft_above    int64  
sqft_basement float64  
yr_built      int64  
zipcode       int64  
lat           float64  
long          float64  
sqft_living15 int64  
sqft_lot15    int64  
renovated     int32  
dtype: object
```

EXPLORE

After cleaning the data, I then explore it to understand its structure, trends, and patterns. This typically involves generating descriptive statistics, and visualizing the data using various kinds of plots. The goal is to gain insights that will help you when modeling.

Identify categorical variables

In [50]:

```
df
```

Out[50]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
0	2014-10-13	221900	3	1.00	1180	5650	1.0	0.0	0.0	0.0
1	2014-12-09	538000	3	2.25	2570	7242	2.0	0.0	0.0	0.0
2	2015-02-25	180000	2	1.00	770	10000	1.0	0.0	0.0	0.0
3	2014-12-09	604000	4	3.00	1960	5000	1.0	0.0	0.0	0.0
4	2015-02-18	510000	3	2.00	1680	8080	1.0	0.0	0.0	0.0
...
21592	2014-05-21	360000	3	2.50	1530	1131	3.0	0.0	0.0	0.0
21593	2015-02-23	400000	4	2.50	2310	5813	2.0	0.0	0.0	0.0
21594	2014-06-23	402101	2	0.75	1020	1350	2.0	0.0	0.0	0.0
21595	2015-01-16	400000	3	2.50	1600	2388	2.0	0.0	0.0	0.0
21596	2014-10-15	325000	2	0.75	1020	1076	2.0	0.0	0.0	0.0

15334 rows × 20 columns

In [51]:

```
#fill out NaN values for Zipcode
df.zipcode.mean()
df.zipcode.median()
df.zipcode.fillna(df.zipcode.median(), inplace=True)
df.zipcode.isnull().sum()
```

Out[51]:

0

In [52]:

df.tail()

Out[52]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
21592	2014-05-21	360000	3	2.50	1530	1131	3.0	0.0	0.0	
21593	2015-02-23	400000	4	2.50	2310	5813	2.0	0.0	0.0	
21594	2014-06-23	402101	2	0.75	1020	1350	2.0	0.0	0.0	
21595	2015-01-16	400000	3	2.50	1600	2388	2.0	0.0	0.0	
21596	2014-10-15	325000	2	0.75	1020	1076	2.0	0.0	0.0	

◀ ▶

In [53]:

df.reset_index(drop = True)

Out[53]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
0	2014-10-13	221900	3	1.00	1180	5650	1.0	0.0	0.0	
1	2014-12-09	538000	3	2.25	2570	7242	2.0	0.0	0.0	
2	2015-02-25	180000	2	1.00	770	10000	1.0	0.0	0.0	
3	2014-12-09	604000	4	3.00	1960	5000	1.0	0.0	0.0	
4	2015-02-18	510000	3	2.00	1680	8080	1.0	0.0	0.0	
...
15329	2014-05-21	360000	3	2.50	1530	1131	3.0	0.0	0.0	
15330	2015-02-23	400000	4	2.50	2310	5813	2.0	0.0	0.0	
15331	2014-06-23	402101	2	0.75	1020	1350	2.0	0.0	0.0	
15332	2015-01-16	400000	3	2.50	1600	2388	2.0	0.0	0.0	
15333	2014-10-15	325000	2	0.75	1020	1076	2.0	0.0	0.0	

15334 rows × 20 columns

◀ ▶

In [54]:

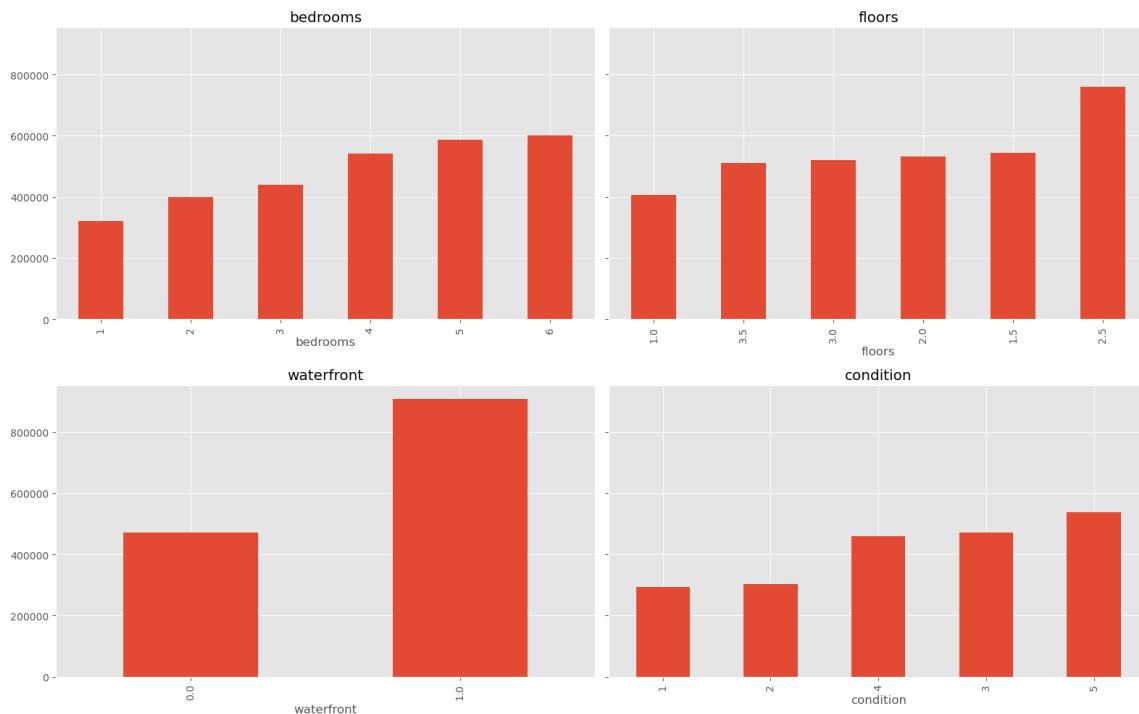
```
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(16,10), sharey=True)

cat = ['bedrooms', 'floors', 'waterfront', 'condition']

for col, ax in zip(cat, axes.flatten()):
    (df.groupby(col) # group values together by column of interest
     .mean(numeric_only=True)['price'] # take the mean of the saleprice for each group
     .sort_values() # sort the groups in ascending order
     .plot
     .bar(ax=ax)) # create a bar graph on the axis

    ax.set_title(col) # Make the title the name of the column

fig.tight_layout()
```



In [55]:

```
from sklearn.preprocessing import OneHotEncoder
import pandas as pd

# creating instance of one-hot-encoder
encoder = OneHotEncoder(drop='first')

# Assume df[cat] contains the categorical columns of your original dataframe
# perform one-hot encoding on cat list
encoder_df = pd.DataFrame(encoder.fit_transform(df[cat]).toarray())

# update column names
cols = []
for x in encoder.get_feature_names_out():
    if 'x0' in x:
        cols.append(x.replace('x0', 'bedrooms'))
    elif 'x1' in x:
        cols.append(x.replace('x1', 'floors'))
    elif 'x2' in x:
        cols.append(x.replace('x2', 'waterfront'))
    elif 'x3' in x:
        cols.append(x.replace('x3', 'condition'))
    else:
        cols.append(x)

# set encoder_df columns equal cols
encoder_df.columns = cols

# merge one-hot encoded columns back with original DataFrame
df_ohe = df.join(encoder_df)

# view final df
print(df_ohe.head())
```

```

date      price    bedrooms  bathrooms  sqft_living  sqft_lot   floors
\\
0 2014-10-13  221900           3       1.00      1180      5650     1.0
1 2014-12-09  538000           3       2.25      2570      7242     2.0
2 2015-02-25  180000           2       1.00       770      10000     1.0
3 2014-12-09  604000           4       3.00      1960      5000     1.0
4 2015-02-18  510000           3       2.00      1680      8080     1.0

waterfront  view  condition  ...  floors_1.5  floors_2.0  floors_2.5  \\
0          0.0   0.0          3  ...        0.0        0.0        0.0
1          0.0   0.0          3  ...        0.0        1.0        0.0
2          0.0   0.0          3  ...        0.0        0.0        0.0
3          0.0   0.0          5  ...        0.0        0.0        0.0
4          0.0   0.0          3  ...        0.0        0.0        0.0

floors_3.0  floors_3.5  waterfront_1.0  condition_2  condition_3  \\
0          0.0        0.0        0.0        0.0        1.0
1          0.0        0.0        0.0        0.0        1.0
2          0.0        0.0        0.0        0.0        1.0
3          0.0        0.0        0.0        0.0        0.0
4          0.0        0.0        0.0        0.0        1.0

condition_4  condition_5
0          0.0        0.0
1          0.0        0.0
2          0.0        0.0
3          0.0        1.0
4          0.0        0.0

```

[5 rows x 35 columns]

In [56]:

```
df_ohe.head()
```

Out[56]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condit
0	2014-10-13	221900	3	1.00	1180	5650	1.0	0.0	0.0	
1	2014-12-09	538000	3	2.25	2570	7242	2.0	0.0	0.0	
2	2015-02-25	180000	2	1.00	770	10000	1.0	0.0	0.0	
3	2014-12-09	604000	4	3.00	1960	5000	1.0	0.0	0.0	
4	2015-02-18	510000	3	2.00	1680	8080	1.0	0.0	0.0	

5 rows x 35 columns

In [57]:

```
new_cols = []
for x in df_ohe:
    new_cols.append(x.replace(".", "_"))
df_ohe.columns = new_cols
df_ohe.head()
```

Out[57]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condit
0	2014-10-13	221900	3	1.00	1180	5650	1.0	0.0	0.0	
1	2014-12-09	538000	3	2.25	2570	7242	2.0	0.0	0.0	
2	2015-02-25	180000	2	1.00	770	10000	1.0	0.0	0.0	
3	2014-12-09	604000	4	3.00	1960	5000	1.0	0.0	0.0	
4	2015-02-18	510000	3	2.00	1680	8080	1.0	0.0	0.0	

5 rows × 35 columns

In [58]:

```
# Drop categorical features to avoid redundancy
df_ohe.drop((cat), axis=1, inplace=True)
```

Converting date to age of the house sold

In [59]:

```
# Change date dtype to datetime
df['date'] = pd.to_datetime(df['date'])

df['date'].min()
```

Out[59]:

Timestamp('2014-05-02 00:00:00')

In [60]:

```
# check dtypes
df.dtypes
```

Out[60]:

date	datetime64[ns]
price	int64
bedrooms	int64
bathrooms	float64
sqft_living	int64
sqft_lot	int64
floors	float64
waterfront	float64
view	float64
condition	int64
grade	int64
sqft_above	int64
sqft_basement	float64
yr_built	int64
zipcode	int64
lat	float64
long	float64
sqft_living15	int64
sqft_lot15	int64
renovated	int32
	dtype: object

In [61]:

```
# Create new column 'age at sold' by subtracting 'date' and 'yr_built'
age_at_sold = df['date'].dt.year - df['yr_built']
df['age_at_sold'] = age_at_sold
df.head()
```

Out[61]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condit
0	2014-10-13	221900		3	1.00	1180	5650	1.0	0.0	0.0
1	2014-12-09	538000		3	2.25	2570	7242	2.0	0.0	0.0
2	2015-02-25	180000		2	1.00	770	10000	1.0	0.0	0.0
3	2014-12-09	604000		4	3.00	1960	5000	1.0	0.0	0.0
4	2015-02-18	510000		3	2.00	1680	8080	1.0	0.0	0.0

5 rows × 21 columns

age_at_sold = New column is now a continuous feature.

In [62]:

```
# Drop columns 'date'  
df.drop('date', axis=1, inplace=True)
```

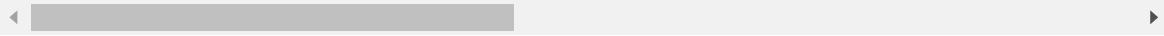
In [63]:

df

Out[63]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
0	221900	3	1.00	1180	5650	1.0	0.0	0.0	3
1	538000	3	2.25	2570	7242	2.0	0.0	0.0	3
2	180000	2	1.00	770	10000	1.0	0.0	0.0	3
3	604000	4	3.00	1960	5000	1.0	0.0	0.0	5
4	510000	3	2.00	1680	8080	1.0	0.0	0.0	3
...
21592	360000	3	2.50	1530	1131	3.0	0.0	0.0	3
21593	400000	4	2.50	2310	5813	2.0	0.0	0.0	3
21594	402101	2	0.75	1020	1350	2.0	0.0	0.0	3
21595	400000	3	2.50	1600	2388	2.0	0.0	0.0	3
21596	325000	2	0.75	1020	1076	2.0	0.0	0.0	3

15334 rows × 20 columns



Check for multicollinearity

Scatter matrix

In [64]:

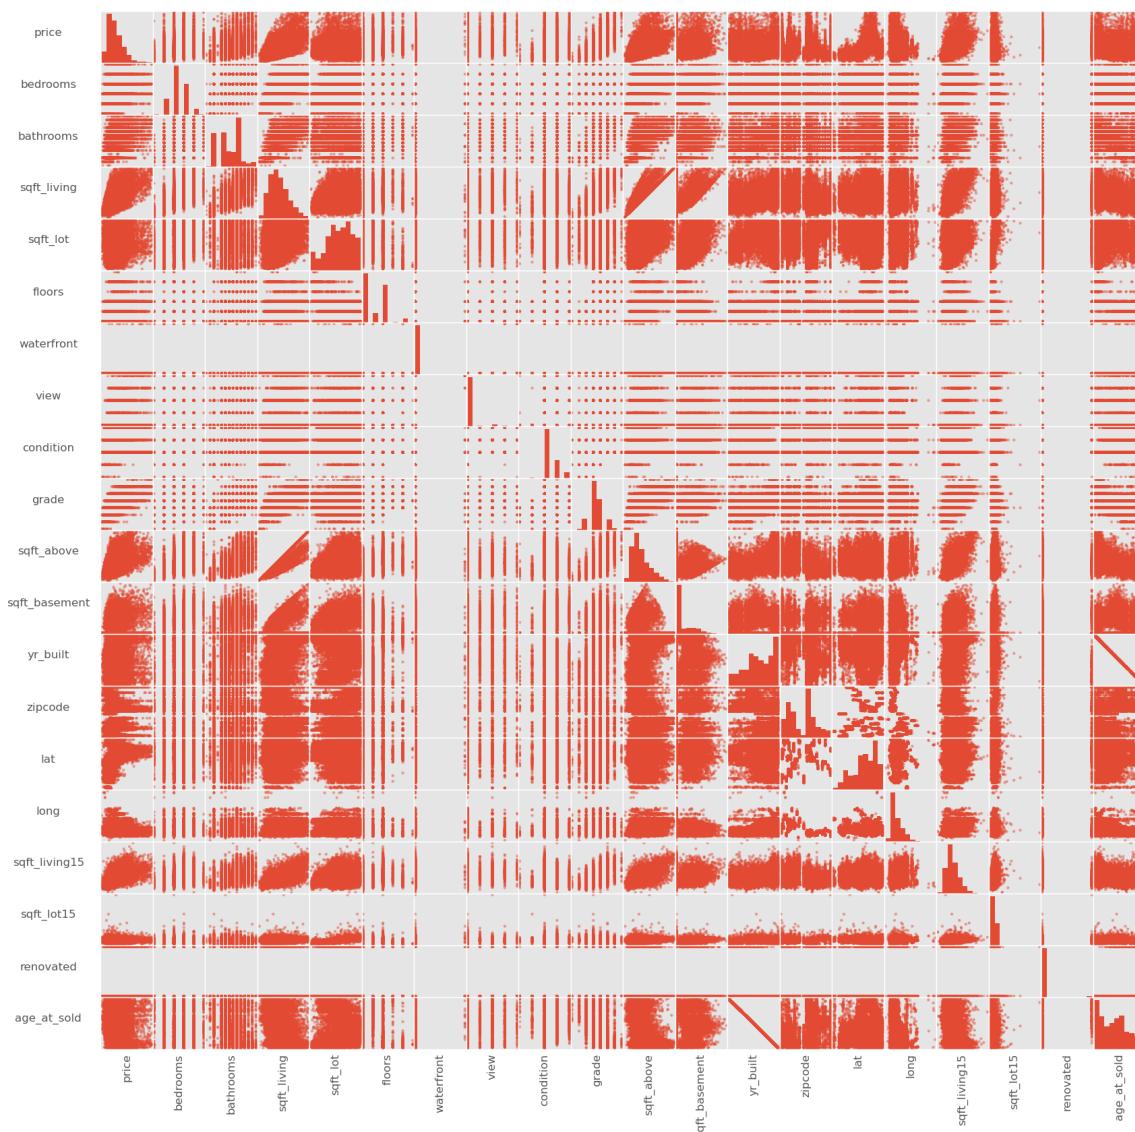
```
sm = pd.plotting.scatter_matrix(df, figsize=[20, 20]);

# Rotates the text
[s.xaxis.label.set_rotation(90) for s in sm.reshape(-1)]
[s.yaxis.label.set_rotation(0) for s in sm.reshape(-1)]

#May need to offset Label when rotating to prevent overlap of figure
[s.get_yaxis().set_label_coords(-1,0.5) for s in sm.reshape(-1)]

#Hide all ticks
[s.set_xticks([]) for s in sm.reshape(-1)]
[s.set_yticks([]) for s in sm.reshape(-1)]

plt.show()
```



Correlation matrix

In [65]:

```
#Correlation matrix
df.corr(numeric_only=True)
```

Out[65]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
price	1.000000	0.270020	0.401291	0.582435	-0.091096	0.258474	0.062988
bedrooms	0.270020	1.000000	0.484538	0.628025	0.230280	0.159391	-0.026551
bathrooms	0.401291	0.484538	1.000000	0.701893	-0.092625	0.533343	-0.009154
sqft_living	0.582435	0.628025	0.701893	1.000000	0.171210	0.343853	-0.006088
sqft_lot	-0.091096	0.230280	-0.092625	0.171210	1.000000	-0.449486	0.007880
floors	0.258474	0.159391	0.533343	0.343853	-0.449486	1.000000	0.003375
waterfront	0.062988	-0.026551	-0.009154	-0.006088	0.007880	0.003375	1.000000
view	0.288403	0.031689	0.082679	0.160663	0.024788	-0.004820	0.200719
condition	0.062857	0.015076	-0.152621	-0.069797	0.143082	-0.284153	0.008267
grade	0.600733	0.325391	0.608525	0.670470	-0.047194	0.487685	-0.006307
sqft_above	0.456314	0.507400	0.626796	0.840791	0.136252	0.537249	-0.007722
sqft_basement	0.269508	0.262739	0.194036	0.365580	0.077200	-0.294297	0.002646
yr_built	-0.033115	0.165549	0.555123	0.328073	-0.089224	0.523722	-0.016596
zipcode	0.058303	-0.169524	-0.225719	-0.192169	-0.245899	-0.083009	0.017435
lat	0.443162	-0.075943	-0.041972	-0.005617	-0.191920	0.022036	-0.005560
long	-0.030962	0.168526	0.263014	0.264305	0.219901	0.150705	-0.005365
sqft_living15	0.483158	0.403860	0.519489	0.728039	0.199125	0.275359	0.010680
sqft_lot15	-0.096889	0.180471	-0.089923	0.132528	0.832437	-0.394623	0.051694
renovated	0.107844	0.003951	0.026586	0.031777	-0.014106	-0.006285	0.029324
age_at_sold	0.033206	-0.165735	-0.555538	-0.328501	0.089285	-0.523947	0.016707

In [66]:

```
#Return True for positive or negative correlations that are bigger than 0.75 in the corre
abs(df.corr(numeric_only=True)) > 0.75
```

Out[66]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_built	zipcode	lat	long	sqft_living15	sqft_lot15	renovated	age_at_sold
price	True	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
bedrooms	False	True	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
bathrooms	False	False	True	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
sqft_living	False	False	False	True	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
sqft_lot	False	False	False	False	True	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
floors	False	False	False	False	False	True	False	False	False	False	False	False	False	False	False	False	False	False	False	
waterfront	False	False	False	False	False	False	False	True	False	False	False	False	False	False	False	False	False	False	False	
view	False	False	False	False	False	False	False	False	False	False	False	False	False	False	True	False	False	False	False	
condition	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
grade	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
sqft_above	False	False	False	True	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
sqft_basement	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
yr_built	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
zipcode	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
lat	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
long	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
sqft_living15	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
sqft_lot15	False	False	False	False	True	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
renovated	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	
age_at_sold	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False	

Now, include stack and zip to create a more robust solution that will return the variable pairs from the correlation matrix that have correlations over .75, but less than 1.

In [67]:

```
df1 = df.corr(numeric_only=True).abs().stack().reset_index().sort_values(0, ascending=False)

df1['pairs'] = list(zip(df1.level_0, df1.level_1))

df1.set_index(['pairs'], inplace = True)

df1.drop(columns=['level_1', 'level_0'], inplace = True)

# cc for correlation coefficient
df1.columns = ['cc']

df1.drop_duplicates(inplace=True)

df1[(df1.cc > .75) & (df1.cc < 1)]
```

Out[67]:

	cc
pairs	
(age_at_sold, yr_built)	0.999890
(sqft_above, sqft_living)	0.840791
(sqft_lot, sqft_lot15)	0.832437

There are three sets of variables that are highly correlated.

- age_at_sold with yr_built
- sqft_above with sqft_living
- sqft_lot15 with sqft_lot).

Did not remove any variables yet

Heatmap

In [68]:

```
# View columns
df.columns
```

Out[68]:

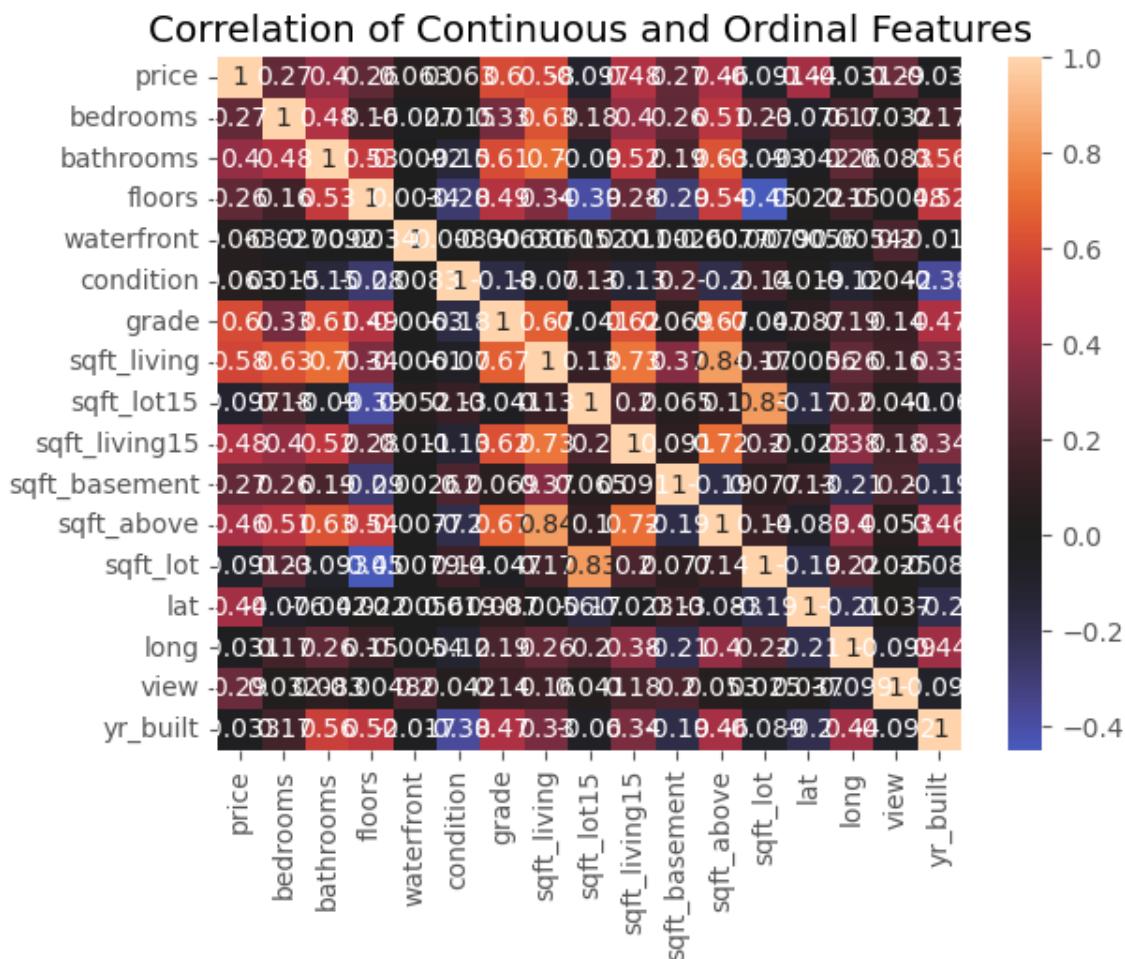
```
Index(['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
       'waterfront', 'view', 'condition', 'grade', 'sqft_above',
       'sqft_basement', 'yr_built', 'zipcode', 'lat', 'long', 'sqft_living15',
       'sqft_lot15', 'renovated', 'age_at_sold'],
      dtype='object')
```

In [69]:

```
# Create a list of continuous and ordinal categorical features
cont_ordinal = ['price', 'bedrooms', 'bathrooms', 'floors', 'waterfront', 'condition', 'grade', 'sqft_living', 'sqft_lot15', 'sqft_living15', 'sqft_basement', 'sqft_above', 'sqft_lot', 'lat', 'long', 'view', 'yr_built']
```

In [70]:

```
# Correlation heat map with continuous and ordinal features
corr = df[cont_ordinal].corr()
sns.heatmap(corr, center=0, annot=True).set_title('Correlation of Continuous and Ordinal Features')
```



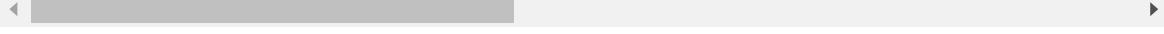
In [71]:

df

Out[71]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
0	221900	3	1.00	1180	5650	1.0	0.0	0.0	3
1	538000	3	2.25	2570	7242	2.0	0.0	0.0	3
2	180000	2	1.00	770	10000	1.0	0.0	0.0	3
3	604000	4	3.00	1960	5000	1.0	0.0	0.0	5
4	510000	3	2.00	1680	8080	1.0	0.0	0.0	3
...
21592	360000	3	2.50	1530	1131	3.0	0.0	0.0	3
21593	400000	4	2.50	2310	5813	2.0	0.0	0.0	3
21594	402101	2	0.75	1020	1350	2.0	0.0	0.0	3
21595	400000	3	2.50	1600	2388	2.0	0.0	0.0	3
21596	325000	2	0.75	1020	1076	2.0	0.0	0.0	3

15334 rows × 20 columns



Can see a correlation between (sqft_above, sqft_living) is 0.840791 and (sqft_lot15, sqft_lot) is 0.83243

Double checking for null values

In [72]:

df

Out[72]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
0	221900	3	1.00	1180	5650	1.0	0.0	0.0	3
1	538000	3	2.25	2570	7242	2.0	0.0	0.0	3
2	180000	2	1.00	770	10000	1.0	0.0	0.0	3
3	604000	4	3.00	1960	5000	1.0	0.0	0.0	5
4	510000	3	2.00	1680	8080	1.0	0.0	0.0	3
...
21592	360000	3	2.50	1530	1131	3.0	0.0	0.0	3
21593	400000	4	2.50	2310	5813	2.0	0.0	0.0	3
21594	402101	2	0.75	1020	1350	2.0	0.0	0.0	3
21595	400000	3	2.50	1600	2388	2.0	0.0	0.0	3
21596	325000	2	0.75	1020	1076	2.0	0.0	0.0	3

15334 rows × 20 columns

In [73]:

```
#check for null values
df.isnull().sum()
```

Out[73]:

price	0
bedrooms	0
bathrooms	0
sqft_living	0
sqft_lot	0
floors	0
waterfront	0
view	43
condition	0
grade	0
sqft_above	0
sqft_basement	0
yr_built	0
zipcode	0
lat	0
long	0
sqft_living15	0
sqft_lot15	0
renovated	0
age_at_sold	0
dtype: int64	

In [74]:

```
#view has 43 null values  
df['view'].value_counts()
```

Out[74]:

```
0.0    14249  
2.0     550  
3.0     221  
1.0     190  
4.0      81  
Name: view, dtype: int64
```

In [75]:

```
df['view'].mean()
```

Out[75]:

```
0.1489111241907004
```

In [76]:

```
df['view'].median()
```

Out[76]:

```
0.0
```

In [77]:

```
df['view'].fillna(df['view'].median(), inplace=True)
```

In [78]:

```
df.view.isnull().sum()
```

Out[78]:

```
0
```

MODEL

The modeling phase involves selecting appropriate algorithms to analyze the data. This phase usually involves training and testing a model.

MODEL 1

In [79]:

df

Out[79]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
0	221900	3	1.00	1180	5650	1.0	0.0	0.0	3
1	538000	3	2.25	2570	7242	2.0	0.0	0.0	3
2	180000	2	1.00	770	10000	1.0	0.0	0.0	3
3	604000	4	3.00	1960	5000	1.0	0.0	0.0	5
4	510000	3	2.00	1680	8080	1.0	0.0	0.0	3
...
21592	360000	3	2.50	1530	1131	3.0	0.0	0.0	3
21593	400000	4	2.50	2310	5813	2.0	0.0	0.0	3
21594	402101	2	0.75	1020	1350	2.0	0.0	0.0	3
21595	400000	3	2.50	1600	2388	2.0	0.0	0.0	3
21596	325000	2	0.75	1020	1076	2.0	0.0	0.0	3

15334 rows × 20 columns

In [80]:

```
train, test = train_test_split(df, random_state = 150)
train.head()
```

Out[80]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
16187	447000	2	1.00	720	7500	1.0	0.0	2.0	3
6767	809000	4	1.75	1790	8372	1.0	0.0	0.0	4
17872	217000	2	1.00	720	4760	1.0	0.0	0.0	5
20242	778983	4	2.50	2490	5647	2.0	0.0	0.0	3
5748	387000	2	2.25	1230	1280	2.0	0.0	0.0	3

In [81]:

```
# Create model function to fit a linear regression that includes rmse
def model1(train, test):

    target = 'price'
    x_cols = list(train.columns)
    x_cols.remove('price')

    predictors = '+'.join(x_cols)
    formula = target + '~' + predictors
    model = ols(formula=formula, data=train).fit()

    # RMSE
    train_err = (mean_squared_error(train['price'], model.predict(train)))**0.5
    test_err = (mean_squared_error(test['price'], model.predict(test)))**0.5

    print("Train RMSE: ", train_err)
    print("Test RMSE: ", test_err, '\n')

return model
```

In [82]:

```
model1(train, test).summary()
```

Train RMSE: 122255.66801937918

Test RMSE: 121587.4592701153

Out[82]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.706
Model:	OLS	Adj. R-squared:	0.706
Method:	Least Squares	F-statistic:	1452.
Date:	Sat, 23 Sep 2023	Prob (F-statistic):	0.00
Time:	19:00:01	Log-Likelihood:	-1.5103e+05

In [82]: **Observations:** 150 **AIC:** 3.0214e+05
Df Residuals: 148 **BIC:** 3.0225e+05
Df Model: 19
 However this can be improved !

Covariance Type: nonrobust
A skewness value of 0.995 is a moderate positive skew, suggesting that there are somewhat more low values and fewer high values, and the "tail" of the distribution stretches more towards the right (higher values).
 coef std err t $P>|t|$ $[0.025$ $0.975]$

Intercept	-4.992e+07	5.56e+06	-8.978	0.000	-6.08e+07	-3.9e+07
-----------	------------	----------	--------	-------	-----------	----------

bedrooms 1.414e+04 1810.270 -7.813 0.000 1.77e+04 1.06e+04
A kurtosis value of 6.506 is greater than 3, indicating a leptokurtic distribution. This suggests that the dataset has more extreme values (outliers) than we would expect from a normal distribution. It implies that the tails of the distribution are relatively heavy, and that data points are more likely to be found far from the mean compared to a normally distributed dataset.
 sqft_living 18.0345 18.503 6.379 0.000 81.766 154.303

sqft_lot	-6.1074	0.923	-6.617	0.000	-7.917	-4.298
----------	---------	-------	--------	-------	--------	--------

floors	7886.2621	3564.781	2.212	0.027	898.684	1.49e+04
--------	-----------	----------	-------	-------	---------	----------

MODEL 2 - Remove features with p_value less than 0.05

In [83]: **grade** 8.953e+04 1926.123 46.485 0.000 8.58e+04 9.33e+04

```
print(train.columns)
sqft_above 8.953e+04 1926.123 46.485 0.000 8.58e+04 9.33e+04
```

```
print(test.columns)
sqft_basement -13.4485 18.286 -0.735 0.462 -49.292 22.395
```

```
Index(['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'yr_built', 'waterfront', 'view', 'condition', 'grade', 'sqft_above', 'sqft_basement', 'zip_code', 'lat', 'long', 'sqft_living', 'sqft_lot', 'renovated', 'age_at_sold'], dtype='object')
```

```
sqft_living15 40.8444 3.442 11.866 0.000 34.097 47.592
```

```
Index(['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'sqft_lot15', 'renovated', 'age_at_sold'], dtype='object')
```

```
sqft_living15 -4.3054 0.761 -5.661 0.000 -5.796 -2.815
```

```
'waterfront', 'view', 'condition', 'grade', 'sqft_above', 'renovated', 'sqft_basement', 'yr_built', 'zip_code', 'lat', 'long', 'sqft_living', 'sqft_lot15', 'renovated', 'age_at_sold'], dtype='object')
```

```
Omnibus: 2292.811 Durbin-Watson: 1.979
```

```
Prob(Omnibus): 0.000 Jarque-Bera (JB): 7785.939
```

```
Skew: 0.995 Prob(JB): 0.00
```

```
Kurtosis: 6.506 Cond. No. 4.80e+08
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [84]:

[2] The condition number is large, 4.8e+08. This might indicate that there are

~~strongly linearly correlated features in your data~~

```
# remove features with p_value less than 0.05
model2 = model1(train, test).summary()
p_table = model2.tables[1]
p_table = pd.DataFrame(p_table.data)
p_table.columns = p_table.iloc[0]
p_table = p_table.drop(0)
p_table = p_table.set_index(p_table.columns[0])
p_table['P>|t|'] = p_table['P>|t|'].astype(float)
sig = list(p_table[p_table['P>|t|'] < 0.05].index)
sig.remove('Intercept')
print("Total # features: ", len(p_table))
print("Total significant features", len(sig))
p_table.head()
```

Train RMSE: 122255.66801937918

Test RMSE: 121587.4592701153

Total # features: 20

Total significant features 17

Out[84]:

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-4.992e+07	5.56e+06	-8.978	0.0	-6.08e+07	-3.9e+07
bedrooms	-1.414e+04	1810.270	-7.813	0.0	-1.77e+04	-1.06e+04
bathrooms	2.156e+04	2917.233	7.389	0.0	1.58e+04	2.73e+04
sqft_living	118.0345	18.503	6.379	0.0	81.766	154.303
sqft_lot	-6.1074	0.923	-6.617	0.0	-7.917	-4.298

In [85]:

```
# Gathering together significant features
train, test = train[sig+['price']], test[sig+['price']]
```

In [86]:

```
# Applying only features with a p_value less than 0.05
model2 = model1(train, test)
model2.summary()
```

Train RMSE: 122284.7267523019

Test RMSE: 121596.23044719924

Out[86]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.706			
Model:	OLS	Adj. R-squared:	0.706			
Method:	Least Squares	F-statistic:	1623.			
Date:	Sat, 23 Sep 2023	Prob (F-statistic):	0.00			
Time:	19:00:02	Log-Likelihood:	-1.5103e+05			
No. Observations:	11500	AIC:	3.021e+05			
Df Residuals:	11482	BIC:	3.022e+05			
Df Model:	17					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	-4.914e+07	5.55e+06	-8.854	0.000	-6e+07	-3.83e+07
bedrooms	-1.415e+04	1810.415	-7.815	0.000	-1.77e+04	-1.06e+04
bathrooms	2.036e+04	2867.004	7.100	0.000	1.47e+04	2.6e+04
sqft_living	110.7889	3.496	31.688	0.000	103.936	117.642
sqft_lot	-5.7416	0.908	-6.322	0.000	-7.522	-3.961
floors	1.189e+04	3066.633	3.876	0.000	5875.209	1.79e+04
waterfront	3.261e+05	3.77e+04	8.653	0.000	2.52e+05	4e+05
view	4.227e+04	2075.824	20.361	0.000	3.82e+04	4.63e+04
condition	2.664e+04	1968.122	13.537	0.000	2.28e+04	3.05e+04
grade	8.989e+04	1919.873	46.823	0.000	8.61e+04	9.37e+04
yr_built	2.031e+04	2445.193	8.306	0.000	1.55e+04	2.51e+04
zipcode	-263.7018	28.731	-9.178	0.000	-320.020	-207.384
lat	5.595e+05	9266.046	60.377	0.000	5.41e+05	5.78e+05
long	-5.612e+04	1.19e+04	-4.699	0.000	-7.95e+04	-3.27e+04
sqft_living15	42.2596	3.382	12.495	0.000	35.630	48.889
sqft_lot15	-4.3441	0.760	-5.712	0.000	-5.835	-2.853
renovated	2.914e+04	6906.844	4.220	0.000	1.56e+04	4.27e+04
age_at_sold	2.246e+04	2444.354	9.190	0.000	1.77e+04	2.73e+04
Omnibus:	2296.870	Durbin-Watson:	1.979			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	7827.477			
Skew:	0.995	Prob(JB):	0.00			
Kurtosis:	6.518	Cond. No.	4.79e+08			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.79e+08. This might indicate that there are strong multicollinearity or other numerical problems.

Check assumptions

Prepare Continuous Data for Modeling: Checking Linearity

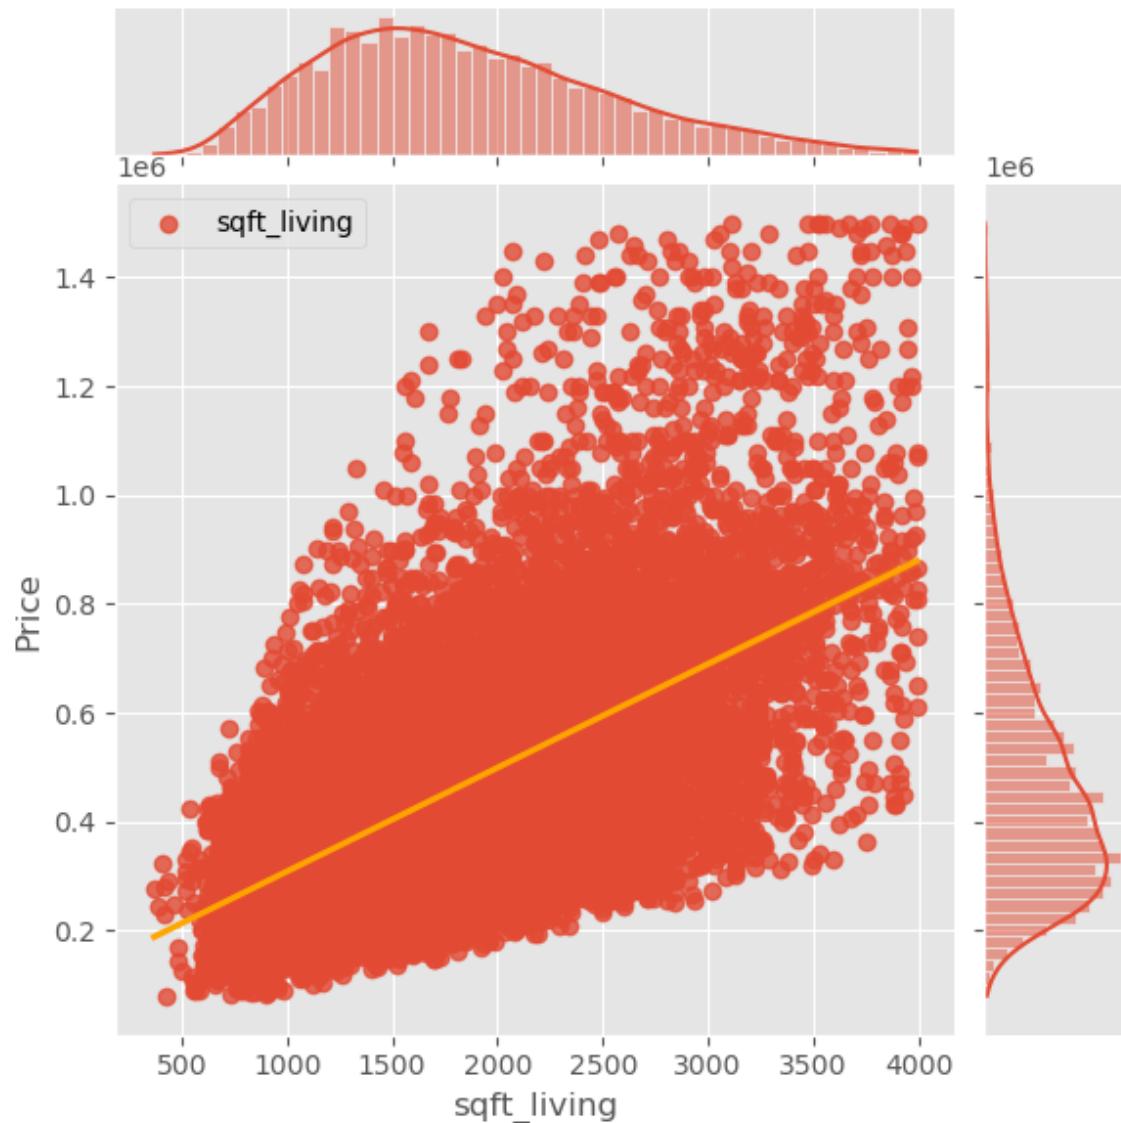
In [87]:

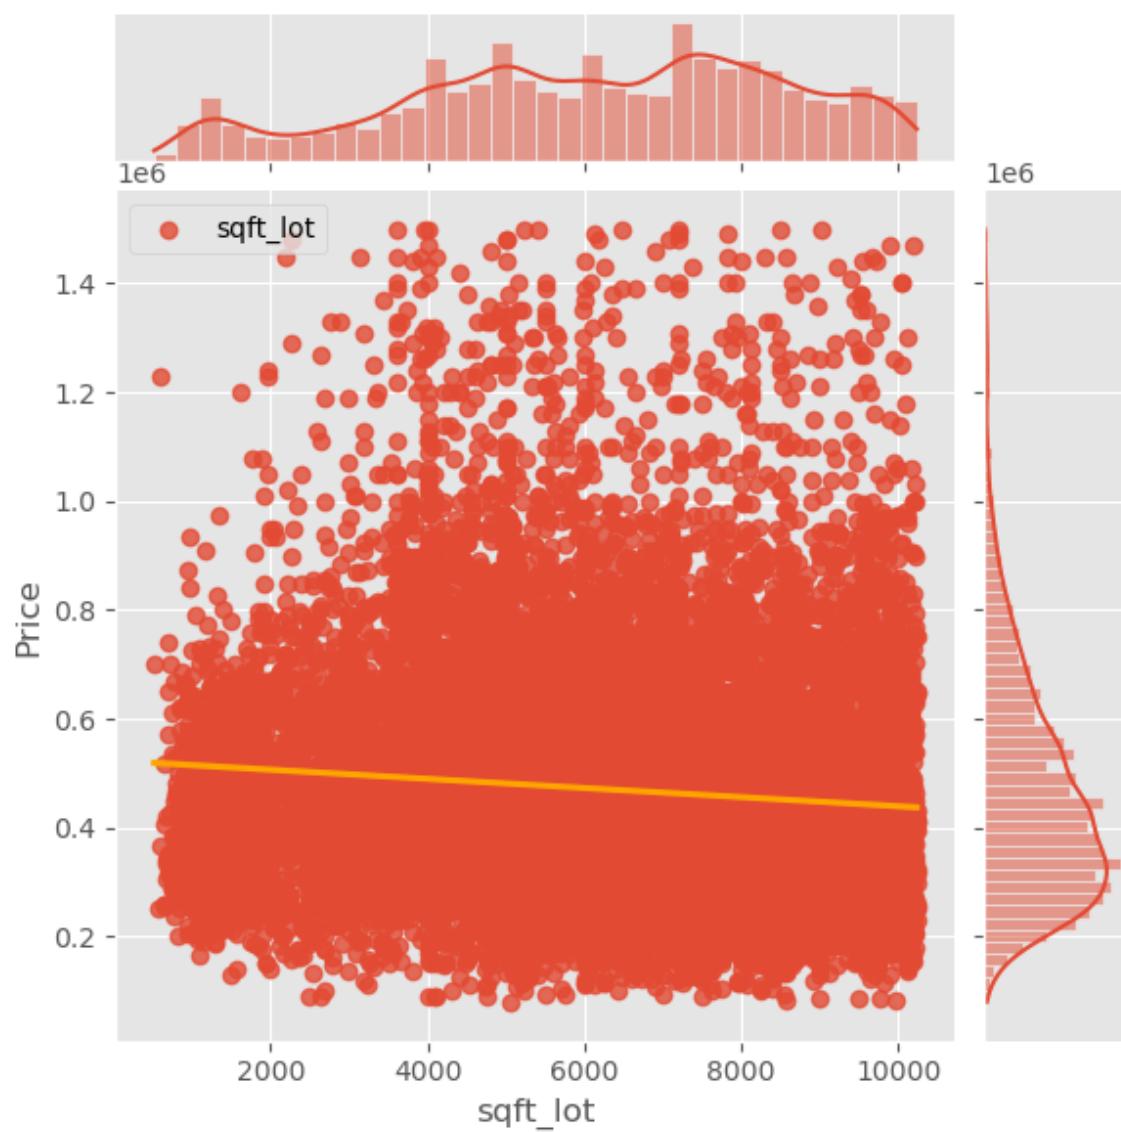
```
# Select and visualize continuous data (data that can take any value)
continuous = ['sqft_living','sqft_lot','sqft_above','sqft_living15','sqft_lot15']

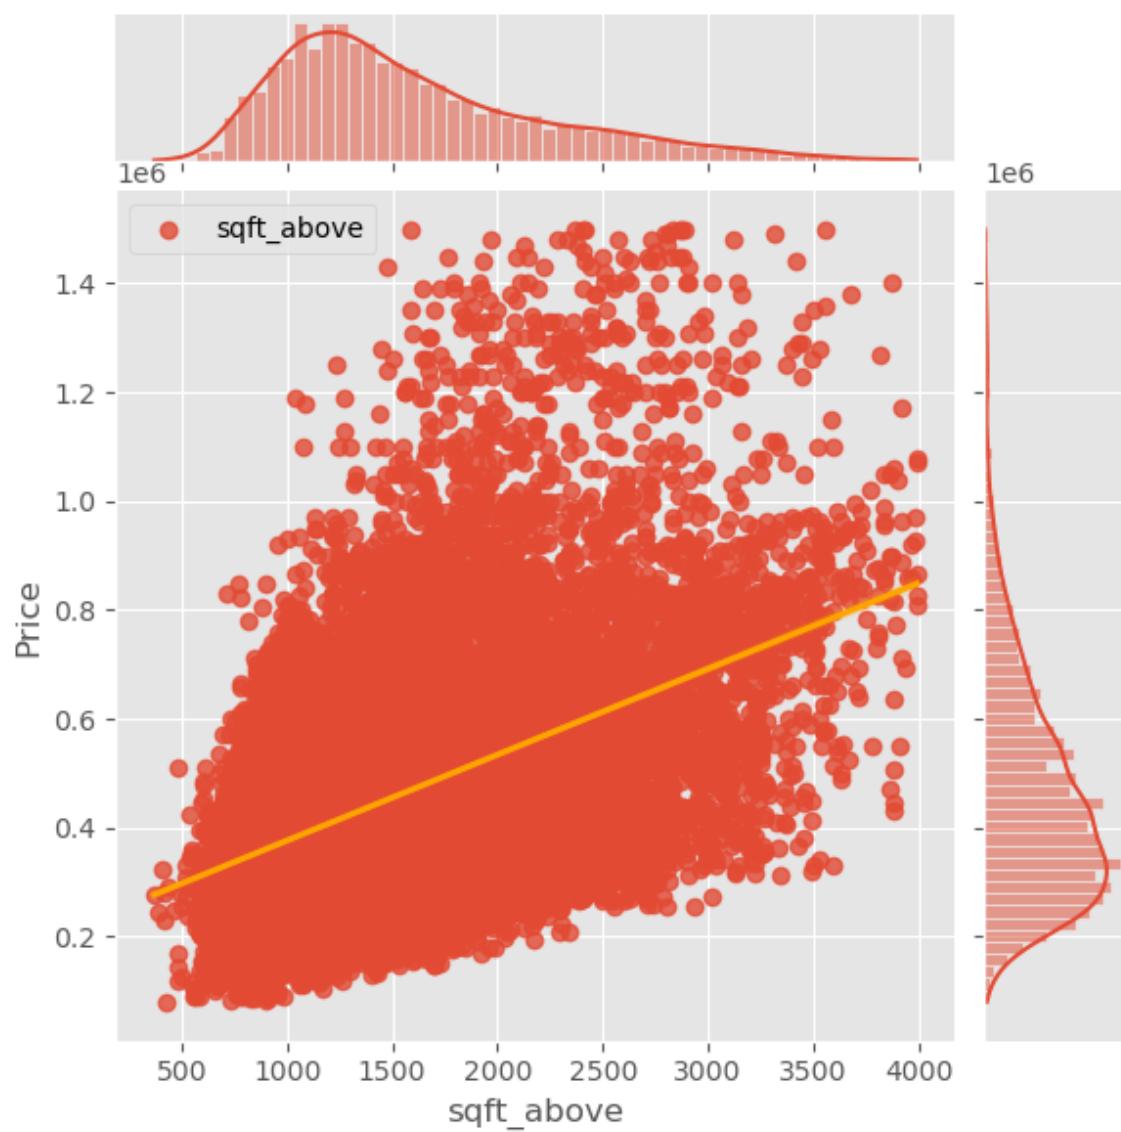
for column in continuous:
    sns.jointplot(x=column, y="price", data=df, kind='reg', label=column,joint_kws={'line_color': 'yellow' })

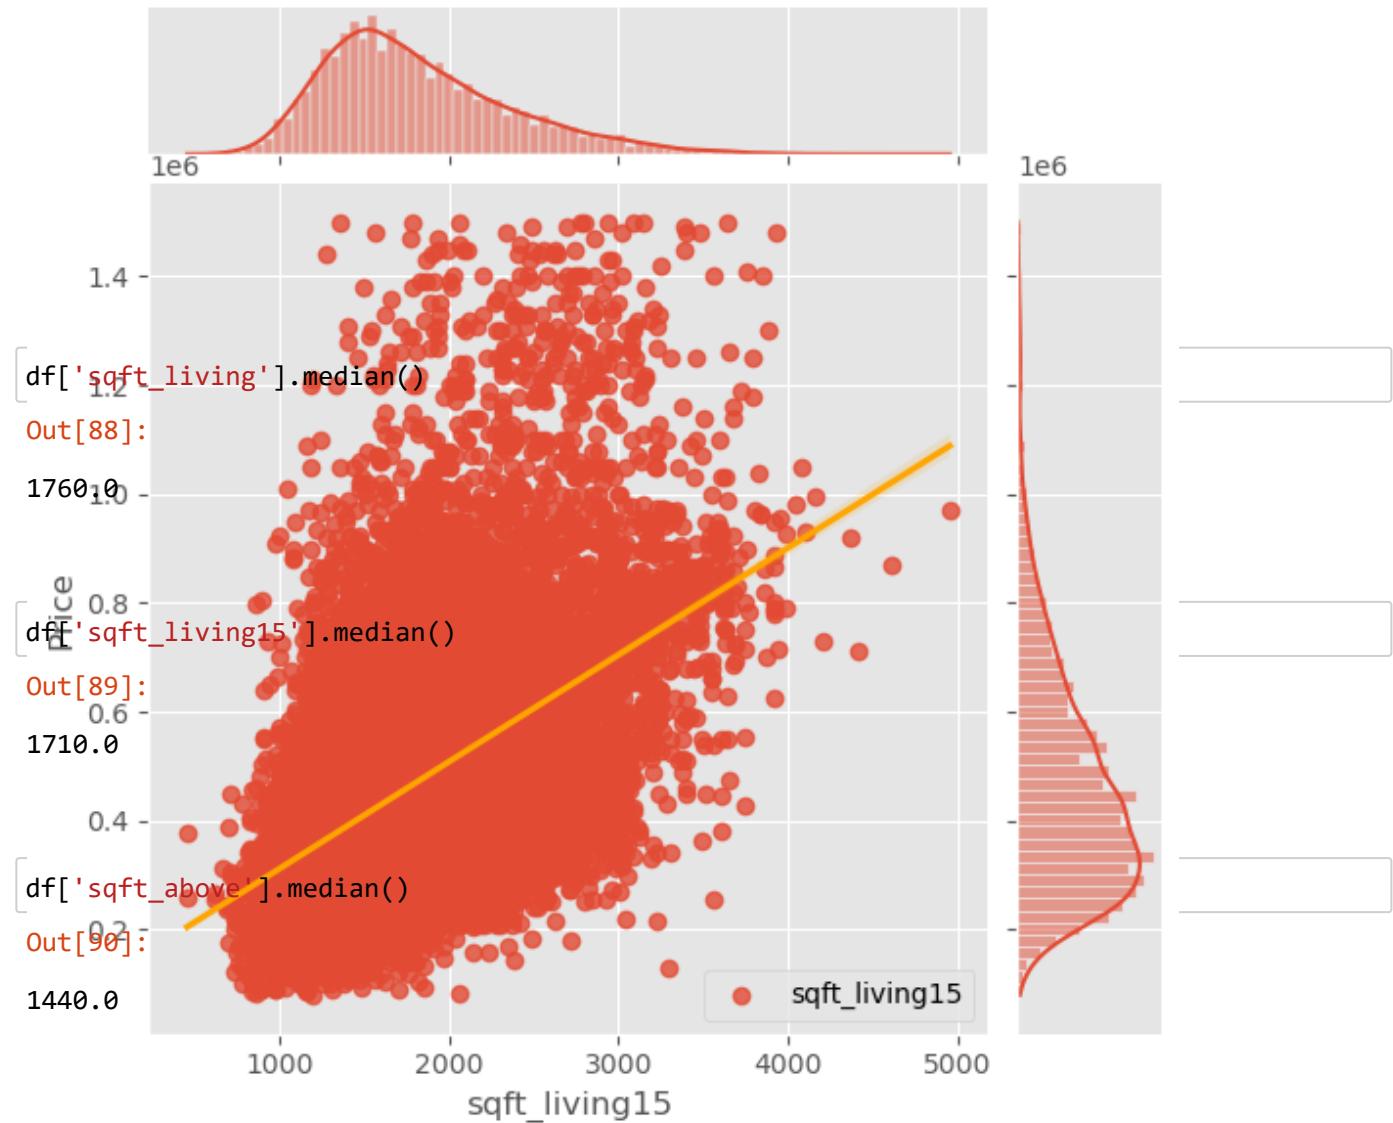
# Set y-axis label
plt.ylabel('Price')

plt.legend()
plt.show()
```



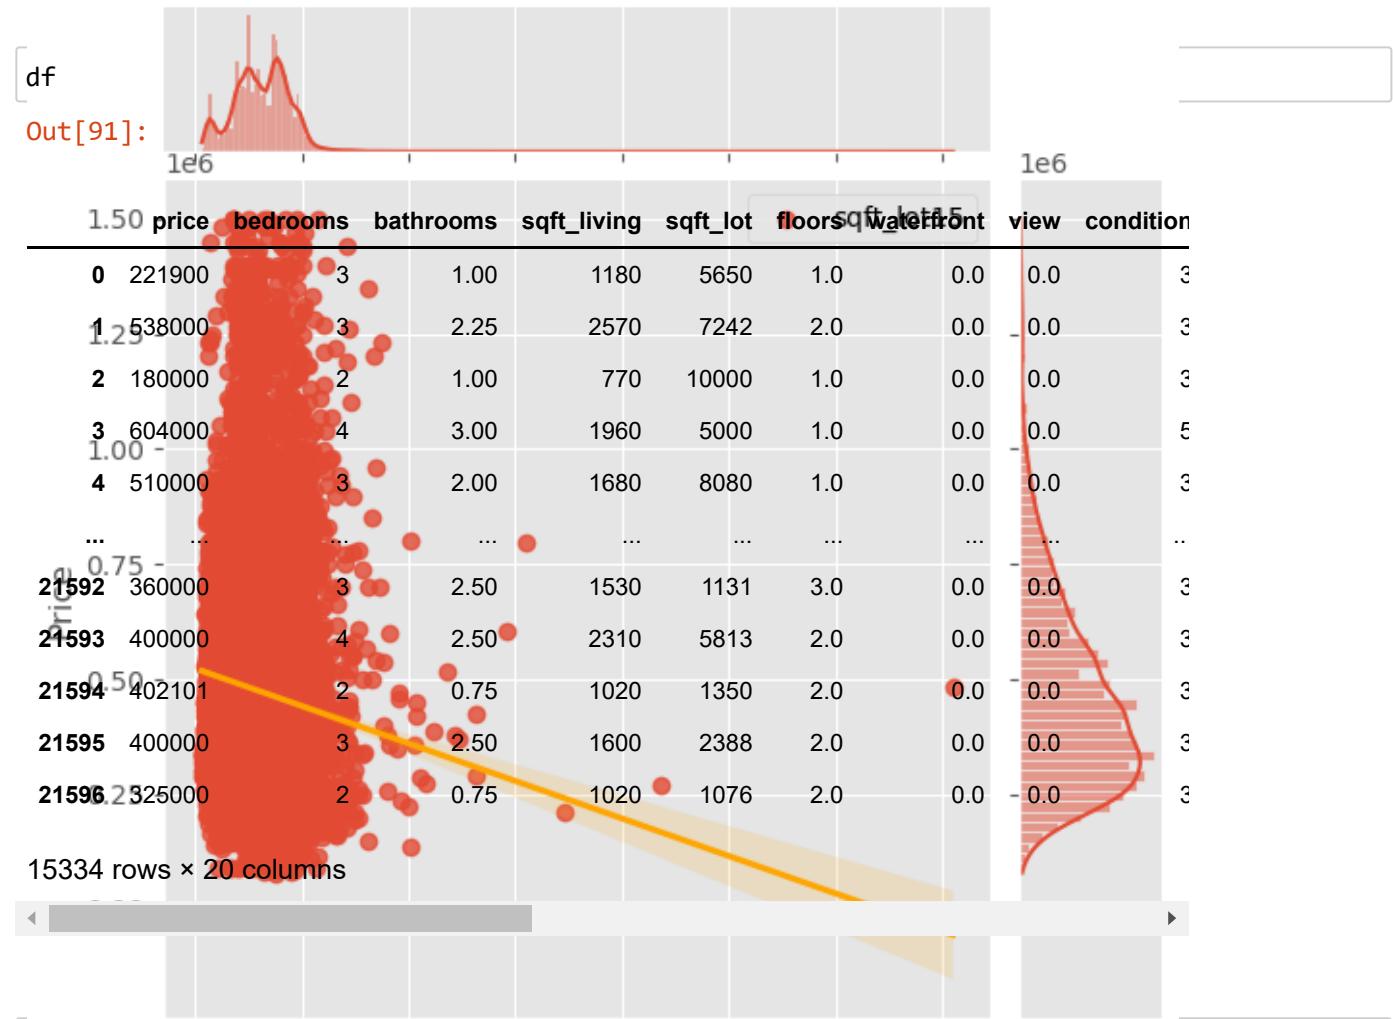






sqft_living, sqft_living15 , sqft_above all have a positive linear relationship with our target variable (price). Since continuous predictors have high multicollinearity (as addressed previously) I will use only sqft_living and sqft_living15

MODEL 3 - addressed multicollinearity



#Removing multicolinear variables (Address the 600 linear)

```
df.drop(['sqft_above', 'sqft_lot15', 'age_at_sold'], axis=1, inplace = True)
df.head()
```

Out[92]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	gr
0	221900	3	1.00	1180	5650	1.0	0.0	0.0	3	
1	538000	3	2.25	2570	7242	2.0	0.0	0.0	3	
2	180000	2	1.00	770	10000	1.0	0.0	0.0	3	
3	604000	4	3.00	1960	5000	1.0	0.0	0.0	5	
4	510000	3	2.00	1680	8080	1.0	0.0	0.0	3	

In [93]:

```
train, test = train_test_split(df, random_state = 150)
train.head()
```

Out[93]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
16187	447000	2	1.00	720	7500	1.0	0.0	2.0	3
6767	809000	4	1.75	1790	8372	1.0	0.0	0.0	4
17872	217000	2	1.00	720	4760	1.0	0.0	0.0	5
20242	778983	4	2.50	2490	5647	2.0	0.0	0.0	3
5748	387000	2	2.25	1230	1280	2.0	0.0	0.0	3

In [94]:

```
# Create model function to fit a Linear regression that includes rmse
def model(train, test):

    target = 'price'
    x_cols = list(train.columns)
    x_cols.remove('price')

    predictors = '+'.join(x_cols)
    formula = target + '~' + predictors
    model = ols(formula=formula, data=train).fit()

    # RMSE
    train_err = (mean_squared_error(train['price'], model.predict(train)))**0.5
    test_err = (mean_squared_error(test['price'], model.predict(test)))**0.5

    print("Train RMSE: ", train_err)
    print("Test RMSE: ", test_err, '\n')

    return model
```

In [95]:

```
model(train, test).summary()
```

Train RMSE: 122874.48068419255

Test RMSE: 122673.49484315024

Out[95]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.703
Model:	OLS	Adj. R-squared:	0.703

After addressing multicollinearity and dropping 'sqft_above', 'sqft_living15', 'age_at_sold' columns there wasn't any improvement in terms of R2 square instead there was a reduction in 0.003

Date: Sat, 23 Sep 2023 Prob(F-statistic): 0.00

A skewness Time of 0.989 suggests a moderately right-skewed distribution. This means:

No. Observations: 11500 AIC: 3.022e+05

- Most of the data points are concentrated on the left side of the distribution.
- Df Residuals: 11483 BIC: 3.023e+05
- The tail of the distribution extends more towards the right, not the left.
- The Df Model: typically greater than the median in a positively skewed distribution.

Covariance Type: nonrobust Kurtosis of 6.416 is leptokurtic, meaning it has fatter tails and a sharper peak compared to a normal

distribution. This would suggest:

	coef	std err	t	P> t	[0.025	0.975]
--	------	---------	---	------	--------	--------

- The data set has a higher propensity for outliers or extreme values than does a normal distribution.
- The distribution is more "peaked" than a normal distribution, implying that data points are more concentrated around the mean.

intercept	4.4240e+04	3.536e+04	1.240	0.000	-4.240e+04	1.510e+05
bedrooms	-1.404e+04	1818.999	-7.719	0.000	-1.76e+04	-1.05e+04
bathrooms	2.194e+04	2926.184	7.498	0.000	1.62e+04	2.77e+04
sqft_living	115.4947	3.902	29.602	0.000	107.847	123.142
sqft_lot	-9.8514	0.641	-15.367	0.000	-11.108	-8.595
floors	8499.5620	3547.739	2.396	0.017	1545.388	1.55e+04
waterfront	3.196e+05	3.79e+04	8.444	0.000	2.45e+05	3.94e+05
view	4.301e+04	2102.288	20.457	0.000	3.89e+04	4.71e+04
condition	2.568e+04	1982.959	12.950	0.000	2.18e+04	2.96e+04
grade	8.942e+04	1934.990	46.210	0.000	8.56e+04	9.32e+04
sqft_basement	-10.7853	4.406	-2.448	0.014	-19.422	-2.149
yr_built	-2178.2131	57.930	-37.600	0.000	-2291.767	-2064.659
zipcode	-262.4633	28.869	-9.091	0.000	-319.052	-205.874
lat	5.596e+05	9353.707	59.826	0.000	5.41e+05	5.78e+05
long	-5.99e+04	1.22e+04	-4.902	0.000	-8.39e+04	-3.59e+04
sqft_living15	38.4202	3.437	11.180	0.000	31.684	45.156
renovated	2.545e+04	6923.065	3.676	0.000	1.19e+04	3.9e+04

Omnibus: 2260.095 Durbin-Watson: 1.979

Prob(Omnibus): 0.000 Jarque-Bera (JB): 7467.245

Skew: 0.989 Prob(JB): 0.00

Kurtosis: 6.416 Cond. No. 2.17e+08

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.17e+08. This might indicate that there are strong multicollinearity or other numerical problems.

NO SIGNIFICANT DIFFERENCES WERE FOUND IN MODEL 3 FROM MODEL 1 AFTER ADDRESSING THE MULTICOLLINEARITY

Log transformation and normalisation only in continuous data

In [96]:

```
df
```

Out[96]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
0	221900	3	1.00	1180	5650	1.0	0.0	0.0	3
1	538000	3	2.25	2570	7242	2.0	0.0	0.0	3
2	180000	2	1.00	770	10000	1.0	0.0	0.0	3
3	604000	4	3.00	1960	5000	1.0	0.0	0.0	5
4	510000	3	2.00	1680	8080	1.0	0.0	0.0	3
...
21592	360000	3	2.50	1530	1131	3.0	0.0	0.0	3
21593	400000	4	2.50	2310	5813	2.0	0.0	0.0	3
21594	402101	2	0.75	1020	1350	2.0	0.0	0.0	3
21595	400000	3	2.50	1600	2388	2.0	0.0	0.0	3
21596	325000	2	0.75	1020	1076	2.0	0.0	0.0	3

15334 rows × 17 columns

In [97]:

```
df.describe()
```

Out[97]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
count	1.533400e+04	15334.000000	15334.000000	15334.000000	15334.000000	15334.000000
mean	4.721766e+05	3.262880	2.007125	1857.591626	6097.929829	1.504695
std	2.240077e+05	0.852576	0.687194	684.015787	2427.995819	0.553630
min	8.000000e+04	1.000000	0.500000	370.000000	520.000000	1.000000
25%	3.090000e+05	3.000000	1.500000	1340.000000	4400.000000	1.000000
50%	4.250000e+05	3.000000	2.000000	1760.000000	6250.000000	1.500000
75%	5.847125e+05	4.000000	2.500000	2290.000000	8007.500000	2.000000
max	1.500000e+06	6.000000	3.750000	3990.000000	10232.000000	3.500000

MODEL 4 - USING LOG TRANSFORMATION FOR CONTINUOUS DATA

In [98]:

```
outcome = 'price'
x_cols = ['sqft_living', 'sqft_lot', 'sqft_living15', 'sqft_basement']
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula=formula, data=df).fit()
model.summary()
```

Out[98]:

OLS Regression Results

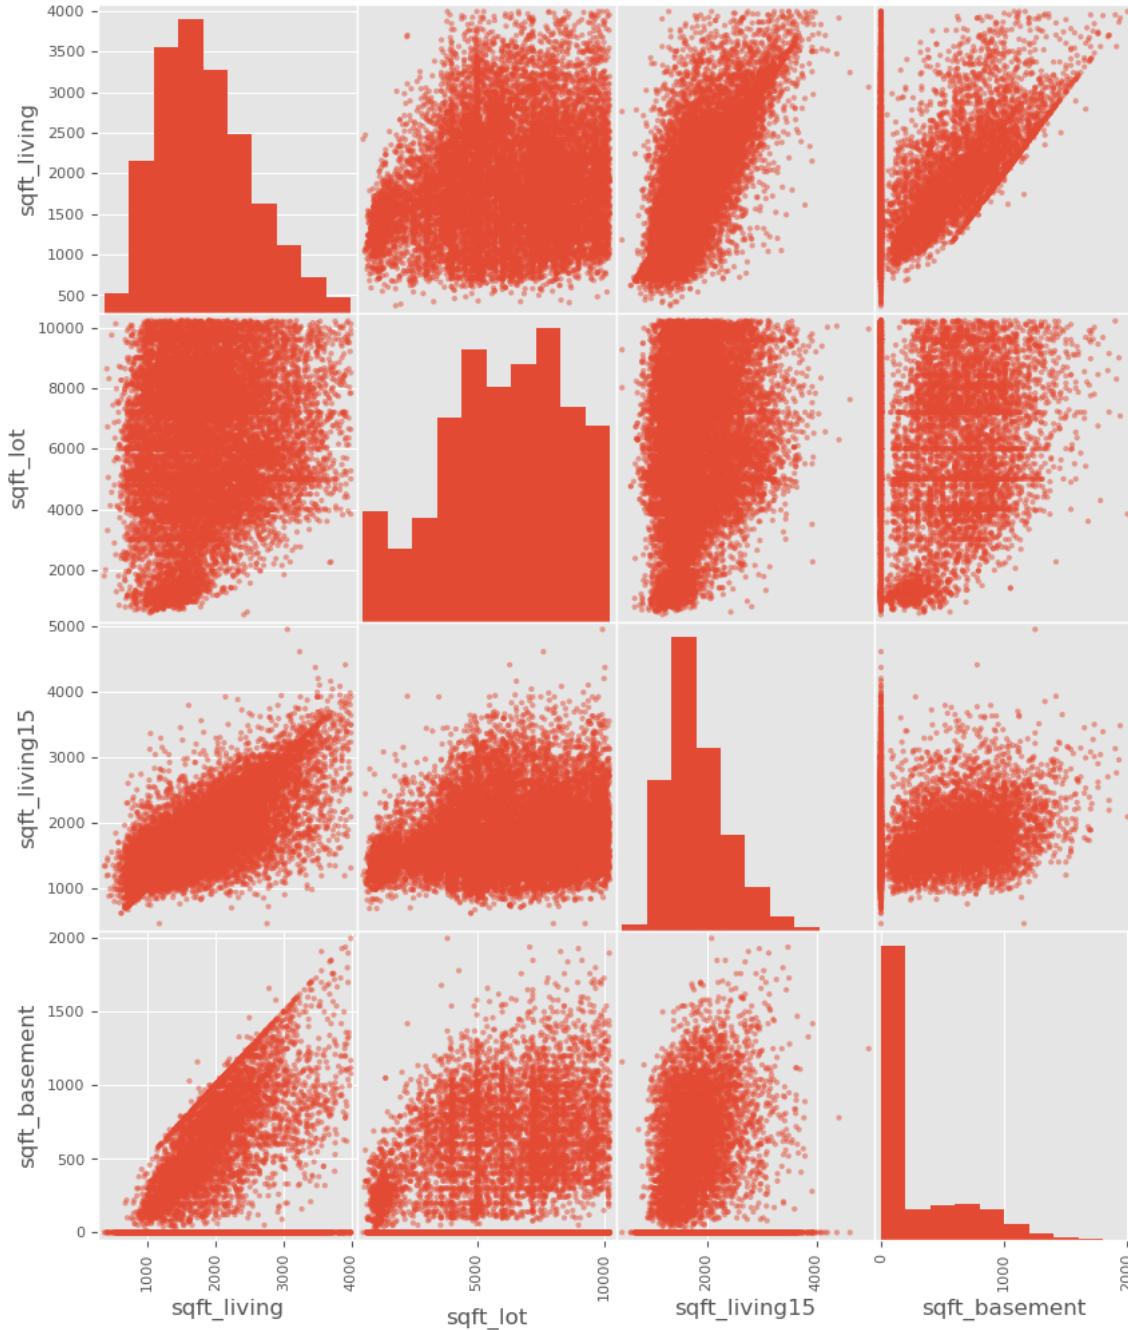
Dep. Variable:	price	R-squared:	0.398			
Model:	OLS	Adj. R-squared:	0.398			
Method:	Least Squares	F-statistic:	2534.			
Date:	Sat, 23 Sep 2023	Prob (F-statistic):	0.00			
Time:	19:00:09	Log-Likelihood:	-2.0677e+05			
No. Observations:	15334	AIC:	4.136e+05			
Df Residuals:	15329	BIC:	4.136e+05			
Df Model:	4					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.644e+05	5554.947	29.595	0.000	1.54e+05	1.75e+05
sqft_living	141.8809	3.330	42.602	0.000	135.353	148.409
sqft_lot	-19.7057	0.591	-33.336	0.000	-20.864	-18.547
sqft_living15	81.6155	3.901	20.922	0.000	73.969	89.262
sqft_basement	65.2759	4.186	15.595	0.000	57.071	73.480
Omnibus:	2414.066	Durbin-Watson:	1.980			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	5066.425			
Skew:	0.946	Prob(JB):	0.00			
Kurtosis:	5.087	Cond. No.	2.78e+04			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.78e+04. This might indicate that there are strong multicollinearity or other numerical problems.

In [99]:

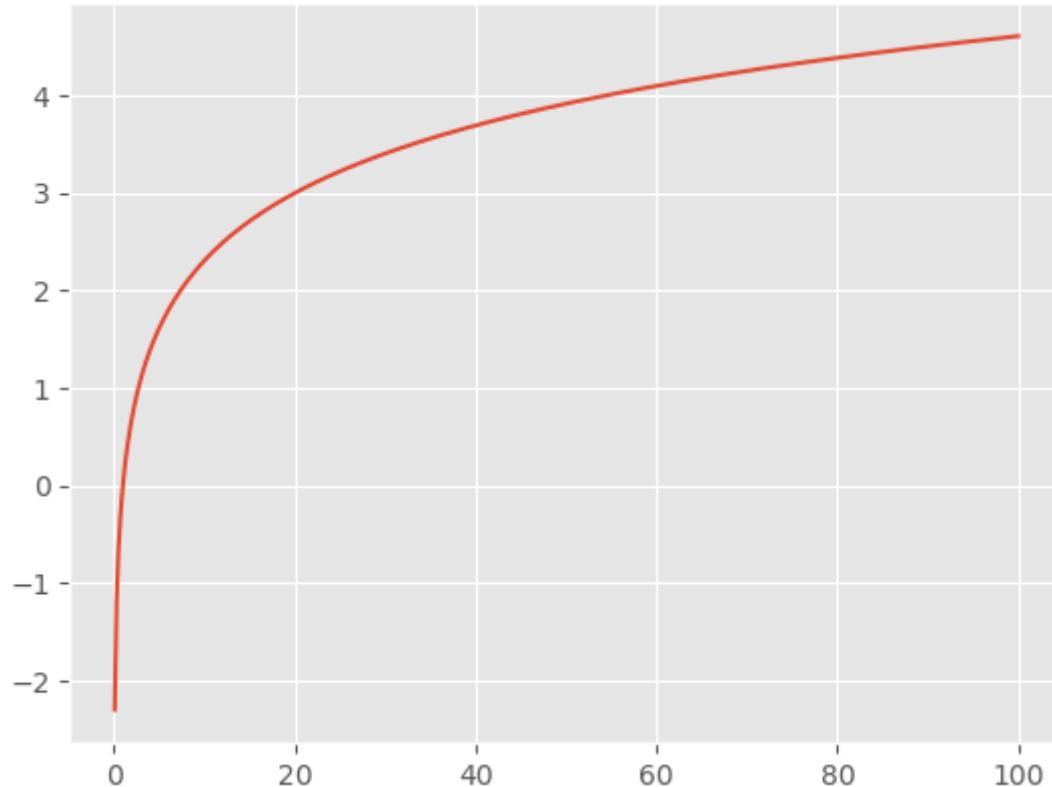
```
pd.plotting.scatter_matrix(df[x_cols], figsize=(10,12));
```



In [100]:

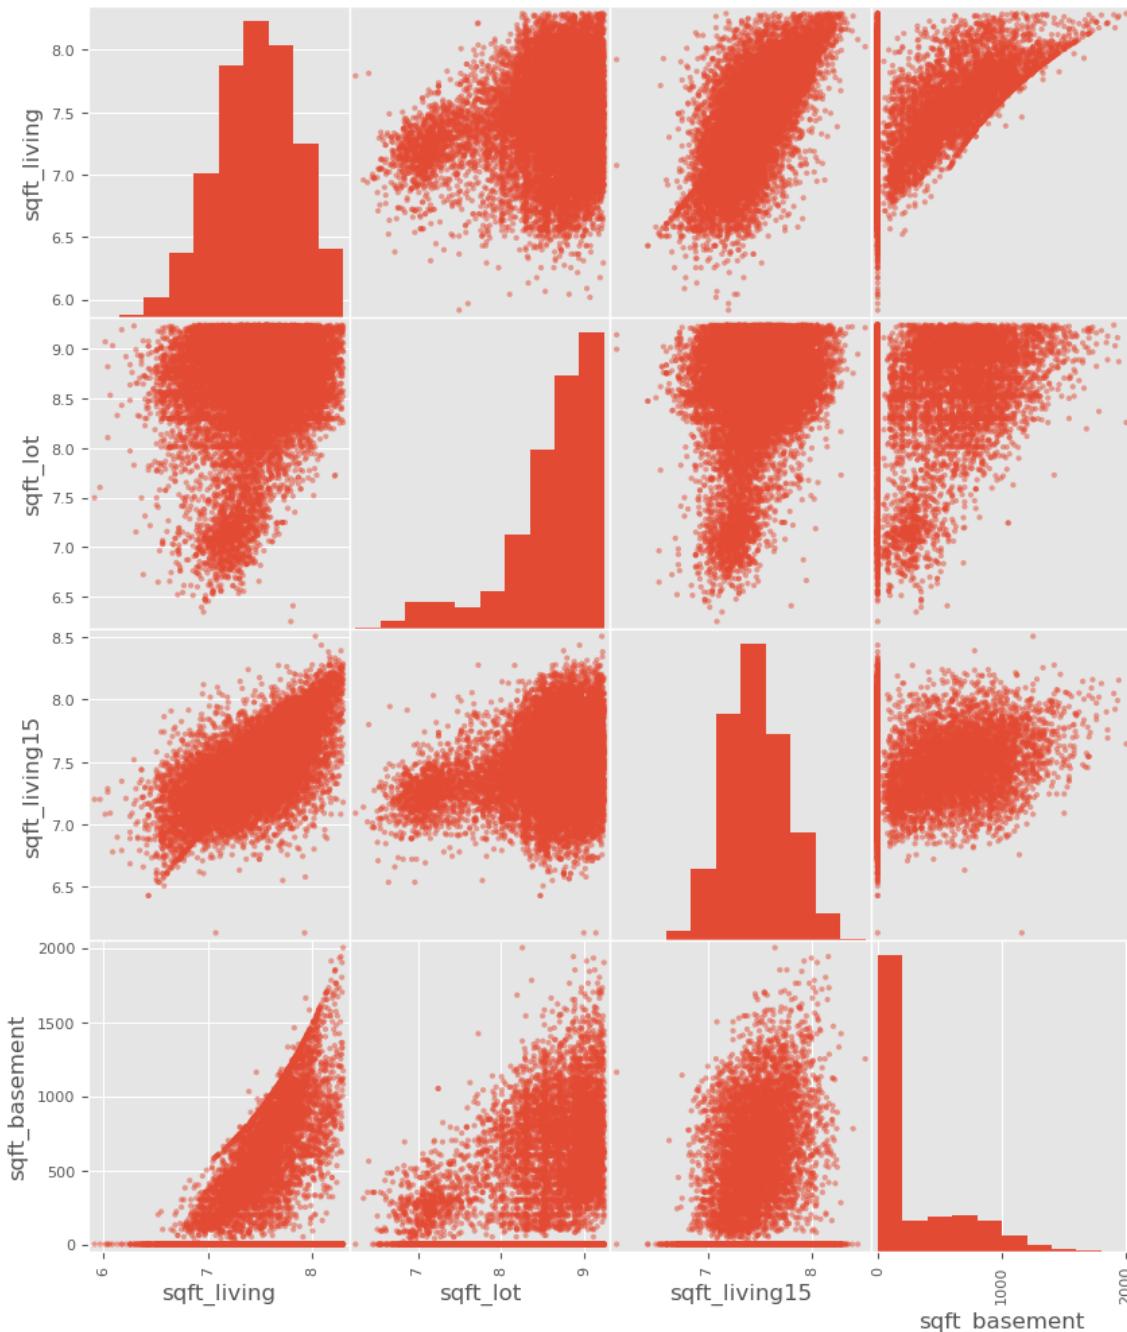
```
x = np.linspace(start=-100, stop=100, num=10**3)
y = np.log(x)
plt.plot(x, y);
```

C:\Users\rani_\AppData\Local\Temp\ipykernel_20364\762412792.py:2: RuntimeWarning: invalid value encountered in log
y = np.log(x)



In [101]:

```
non_normal = ['sqft_living', 'sqft_lot', 'sqft_living15']
for feat in non_normal:
    df[feat] = df[feat].map(lambda x: np.log(x))
pd.plotting.scatter_matrix(df[x_cols], figsize=(10,12));
```



In [102]:

```
outcome = 'price'
x_cols = ['sqft_living', 'sqft_lot', 'sqft_living15', 'sqft_basement']
predictors = '+' .join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula=formula, data=df).fit()
model.summary()
```

Out[102]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.362			
Model:	OLS	Adj. R-squared:	0.361			
Method:	Least Squares	F-statistic:	2170.			
Date:	Sat, 23 Sep 2023	Prob (F-statistic):	0.00			
Time:	19:00:15	Log-Likelihood:	-2.0722e+05			
No. Observations:	15334	AIC:	4.145e+05			
Df Residuals:	15329	BIC:	4.145e+05			
Df Model:	4					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	-1.884e+06	4.02e+04	-46.863	0.000	-1.96e+06	-1.8e+06
sqft_living	2.24e+05	5952.604	37.623	0.000	2.12e+05	2.36e+05
sqft_lot	-7.83e+04	2710.864	-28.885	0.000	-8.36e+04	-7.3e+04
sqft_living15	1.798e+05	7254.859	24.789	0.000	1.66e+05	1.94e+05
sqft_basement	68.8601	4.277	16.099	0.000	60.476	77.244
Omnibus:	2590.857	Durbin-Watson:	1.986			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	5385.142			
Skew:	1.011	Prob(JB):	0.00			
Kurtosis:	5.083	Cond. No.	1.25e+04			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.25e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Comparing OLS regression results before and after log transformation for continuous data

- The R-squared is 0.398 for the first model and 0.362 for the second, which means the model explains around 47-45% of the variability in price.
- The variable sqft_basement seems to be not statistically significant in both models (high p-value).
- Both models suffer from non-normality in the residuals, as evidenced by the Omnibus and JB tests.
- The condition number (Cond. No.) is lower in the second model, which might imply less multicollinearity.

- The increase in skewness from 0.946 to 1.011 indicates that the distribution has become more positively skewed, suggesting more frequent higher values in the dataset.
- The increase in kurtosis from 5.467 to 5.083 suggests that the distribution has become more "tailed,"

INTERPRET

The final step is interpreting the results of the models.

Summary of Models

- Model 1 -> R- SQUARE = 0.706 , Number of Features = 19 , Best Model = No change in data except dropped one column (ID)
- Model 2 -> R- SQUARE = 0.706 , Number of Features = 17 , Best Model = Removed P-value greater than 0.05 for 3 column (sqft_above and sqft_basement)
- Model 3 -> R- SQUARE = 0.703 , Number of Features = 17 , Best Model = Addressed multicollinearity and compared to Model 1
- Model 4 -> R- SQUARE = 0.362 , Number of Features = 5 , Best Model = Applied Log transformation only for continuous data

ANOTHER APPROACH !

Let's first Standardize the variables

MODEL 5 Using selected variables

In [103]:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
col_df = df['sqft_living15'].values.reshape(-1,1)
scale_df = scaler.fit_transform(col_df)
df['sqft_living15_sc'] = scale_df.flatten()

x = df['sqft_living15_sc']
ax = sns.distplot(x)
```

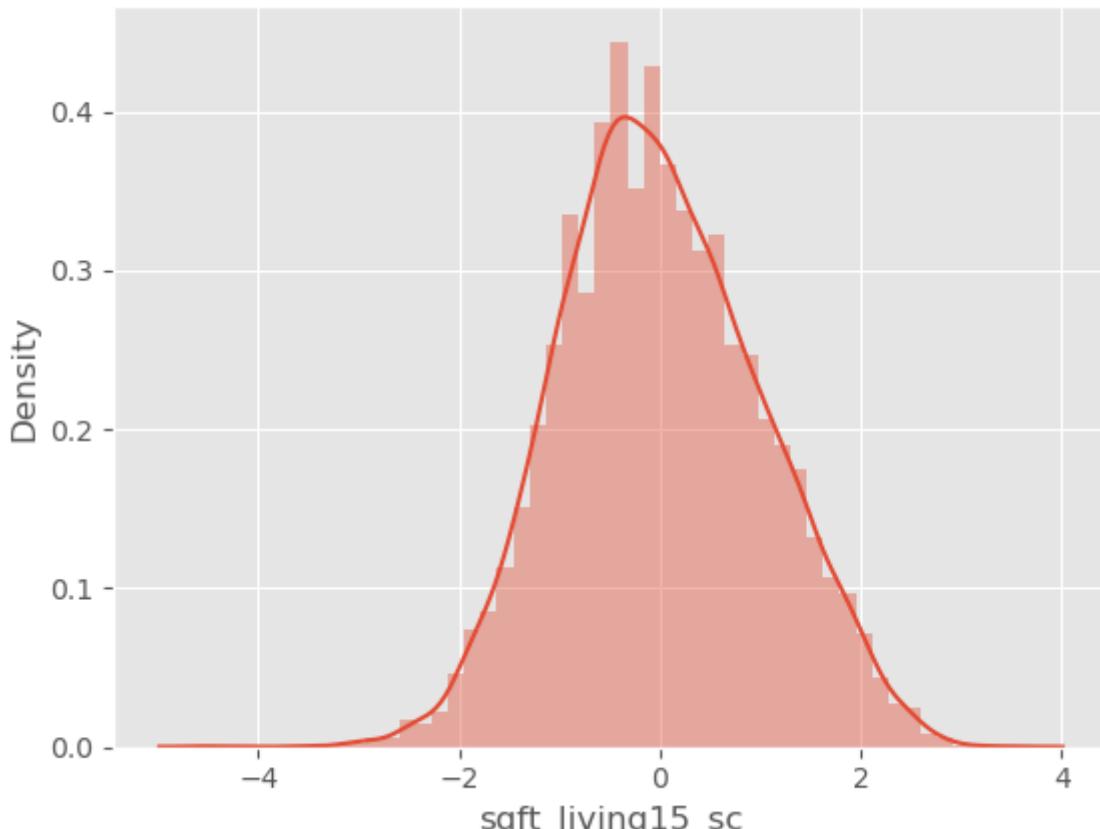
C:\Users\rani_\AppData\Local\Temp\ipykernel_20364\1073268049.py:8: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

```
ax = sns.distplot(x)
```



In [104]:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
col_df = df['sqft_living'].values.reshape(-1,1)
scale_df = scaler.fit_transform(col_df)
df['sqft_living_sc'] = scale_df.flatten()

x = df['sqft_living_sc']
ax = sns.distplot(x)
```

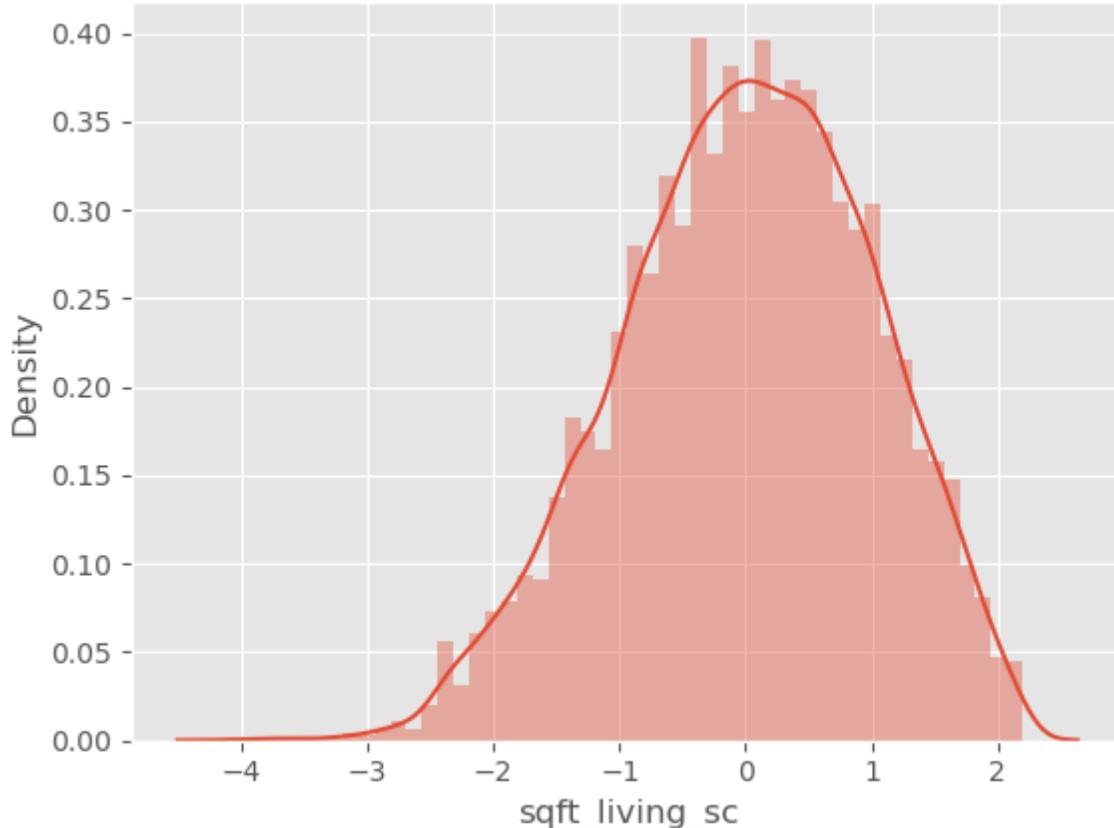
C:\Users\rani_\AppData\Local\Temp\ipykernel_20364\896630809.py:8: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

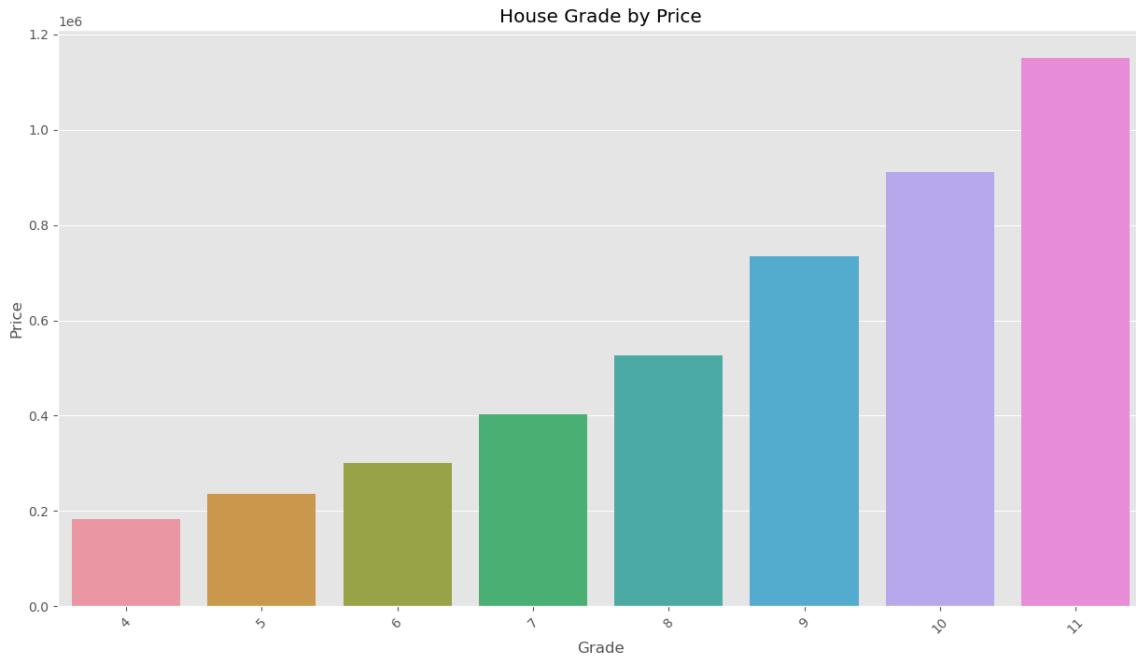
For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

```
ax = sns.distplot(x)
```



In [105]:

```
plt.figure(figsize=(15, 8))
sns.barplot(x='grade', y='price', data=df, errorbar=None)
plt.xticks(rotation=45)
plt.xlabel('Grade')
plt.ylabel('Price')
plt.title('House Grade by Price')
plt.show()
```



Building grade:

Represents the construction quality of improvements. Grades run from grade 1 to 13. Generally defined as:

- 1-3 Falls short of minimum building standards. Normally cabin or inferior structure.
- 4 Generally older, low quality construction. Does not meet code.
- 5 Low construction costs and workmanship. Small, simple design.
- 6 Lowest grade currently meeting building code. Low quality materials and simple designs.
- 7 Average grade of construction and design. Commonly seen in plats and older sub-divisions.
- 8 Just above average in construction and design. Usually better materials in both the exterior and interior finish work.
- 9 Better architectural design with extra interior and exterior design and quality.
- 10 Homes of this quality generally have high quality features. Finish work is better and more design quality is seen in the floor plans. Generally have a larger square footage.
- 11 Custom design and higher quality finish work with added amenities of solid woods, bathroom fixtures and more luxurious options.
- 12 Custom design and excellent builders. All materials are of the highest quality and all conveniences are present.

13 Generally custom designed and built. Mansion level. Large amount of highest quality cabinet work, wood

The visualization above shows that when the building grade improves, the house price rises as well. We can see in the bar plot that the mean house price for a home with a grade of 11 is far above other grades. While the building grade of 4 falls short of minimum building standards based on the King County grading system.

In [106]:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
col_df = df['grade'].values.reshape(-1,1)
scale_df = scaler.fit_transform(col_df)
df['grade_sc'] = scale_df.flatten()

x = df['grade_sc']
ax = sns.distplot(x)
```

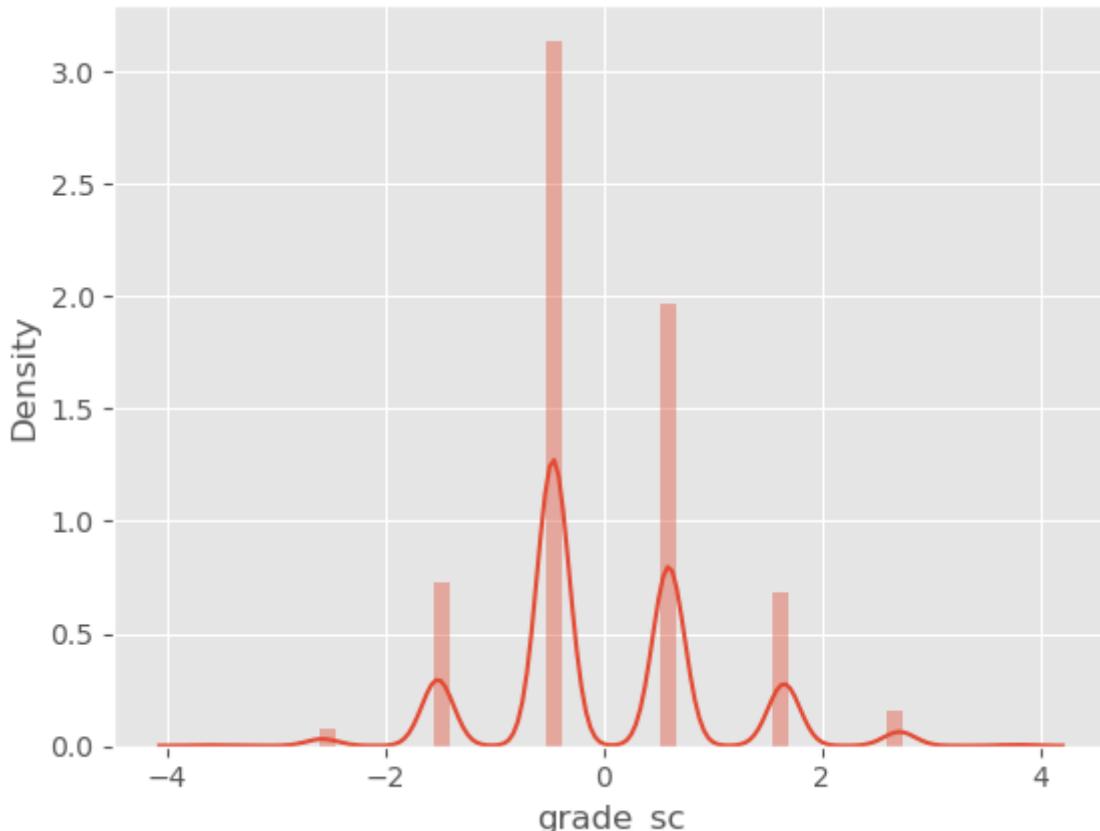
C:\Users\rani_\AppData\Local\Temp\ipykernel_20364\3865672094.py:8: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

```
ax = sns.distplot(x)
```



In [107]:

```
f = 'price~grade_sc+sqft_living_sc+sqft_living15_sc'
model = ols(formula=f, data=df).fit()
model.summary()
```

Out[107]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.404			
Model:	OLS	Adj. R-squared:	0.404			
Method:	Least Squares	F-statistic:	3460.			
Date:	Sat, 23 Sep 2023	Prob (F-statistic):	0.00			
Time:	19:00:16	Log-Likelihood:	-2.0670e+05			
No. Observations:	15334	AIC:	4.134e+05			
Df Residuals:	15330	BIC:	4.134e+05			
Df Model:	3					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	4.722e+05	1397.045	337.982	0.000	4.69e+05	4.75e+05
grade_sc	9.106e+04	1934.308	47.075	0.000	8.73e+04	9.48e+04
sqft_living_sc	5.414e+04	2178.768	24.848	0.000	4.99e+04	5.84e+04
sqft_living15_sc	1.245e+04	2050.301	6.073	0.000	8432.070	1.65e+04
Omnibus:	2142.294	Durbin-Watson:	1.961			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	3825.446			
Skew:	0.914	Prob(JB):	0.00			
Kurtosis:	4.628	Cond. No.	2.88			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We can see that after standardization of our selected features our model performed poorly than before COMPARING TO MODEL 1 AND 2.

In [108]:

```
#include the zipcode column to the model
#zipcodes have no numerical significance in our data, this variable is treated as categorical
f = 'price~grade_sc+sqft_living_sc+sqft_living15_sc+C(zipcode)'
model = ols(formula=f, data=df).fit()
model.summary()
```

Out[108]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.775			
Model:	OLS	Adj. R-squared:	0.774			
Method:	Least Squares	F-statistic:	730.7			
Date:	Sat, 23 Sep 2023	Prob (F-statistic):	0.00			
Time:	19:00:16	Log-Likelihood:	-1.9922e+05			
No. Observations:	15334	AIC:	3.986e+05			
Df Residuals:	15261	BIC:	3.991e+05			
Df Model:	72					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
Intercept	2.625e+05	7103.962	36.951	0.000	2.49e+05	2.76e+05
C(zipcode)[T.98002]	5.112e+04	1.09e+04	4.684	0.000	2.97e+04	7.25e+04
C(zipcode)[T.98003]	1.568e+04	1.05e+04	1.496	0.135	-4859.404	3.62e+04
C(zipcode)[T.98004]	5.824e+05	1.3e+04	44.957	0.000	5.57e+05	6.08e+05
C(zipcode)[T.98005]	3.531e+05	1.68e+04	21.022	0.000	3.2e+05	3.86e+05
C(zipcode)[T.98006]	3.047e+05	1.01e+04	30.143	0.000	2.85e+05	3.25e+05
C(zipcode)[T.98007]	2.585e+05	1.31e+04	19.750	0.000	2.33e+05	2.84e+05
C(zipcode)[T.98008]	2.603e+05	1.03e+04	25.320	0.000	2.4e+05	2.8e+05
C(zipcode)[T.98010]	1.201e+05	2.1e+04	5.713	0.000	7.89e+04	1.61e+05
C(zipcode)[T.98011]	1.485e+05	1.16e+04	12.824	0.000	1.26e+05	1.71e+05
C(zipcode)[T.98014]	1.754e+05	2.61e+04	6.716	0.000	1.24e+05	2.27e+05
C(zipcode)[T.98019]	8.871e+04	1.32e+04	6.704	0.000	6.28e+04	1.15e+05
C(zipcode)[T.98022]	4.537e+04	1.21e+04	3.756	0.000	2.17e+04	6.9e+04
C(zipcode)[T.98023]	-6358.3963	8958.869	-0.710	0.478	-2.39e+04	1.12e+04
C(zipcode)[T.98024]	2.03e+05	3.44e+04	5.894	0.000	1.35e+05	2.7e+05
C(zipcode)[T.98027]	2.324e+05	1.04e+04	22.337	0.000	2.12e+05	2.53e+05
C(zipcode)[T.98028]	1.328e+05	1.1e+04	12.094	0.000	1.11e+05	1.54e+05
C(zipcode)[T.98029]	2.22e+05	9600.360	23.121	0.000	2.03e+05	2.41e+05
C(zipcode)[T.98030]	6114.1759	1.03e+04	0.591	0.555	-1.42e+04	2.64e+04
C(zipcode)[T.98031]	1.473e+04	1.02e+04	1.444	0.149	-5267.858	3.47e+04
C(zipcode)[T.98032]	2.38e+04	1.31e+04	1.812	0.070	-1942.272	4.95e+04
C(zipcode)[T.98033]	3.812e+05	9588.995	39.753	0.000	3.62e+05	4e+05
C(zipcode)[T.98034]	1.966e+05	8827.053	22.277	0.000	1.79e+05	2.14e+05
C(zipcode)[T.98038]	3.113e+04	8686.068	3.584	0.000	1.41e+04	4.82e+04
C(zipcode)[T.98039]	7.727e+05	4.09e+04	18.900	0.000	6.93e+05	8.53e+05
C(zipcode)[T.98040]	4.959e+05	1.37e+04	36.144	0.000	4.69e+05	5.23e+05
C(zipcode)[T.98042]	2.313e+04	9058.715	2.553	0.011	5372.175	4.09e+04

C(zipcode)[T.98045]	1.039e+05	1.47e+04	7.090	0.000	7.52e+04	1.33e+05
C(zipcode)[T.98052]	2.617e+05	8974.764	29.158	0.000	2.44e+05	2.79e+05
C(zipcode)[T.98053]	2.593e+05	1.02e+04	25.526	0.000	2.39e+05	2.79e+05
C(zipcode)[T.98055]	6.24e+04	1.01e+04	6.187	0.000	4.26e+04	8.22e+04
C(zipcode)[T.98056]	1.284e+05	9293.634	13.814	0.000	1.1e+05	1.47e+05
C(zipcode)[T.98058]	4.18e+04	9278.424	4.506	0.000	2.36e+04	6e+04
C(zipcode)[T.98059]	9.91e+04	9680.914	10.237	0.000	8.01e+04	1.18e+05
C(zipcode)[T.98065]	1.366e+05	9873.457	13.839	0.000	1.17e+05	1.56e+05
C(zipcode)[T.98070]	2.178e+05	3.83e+04	5.684	0.000	1.43e+05	2.93e+05
C(zipcode)[T.98072]	1.553e+05	1.32e+04	11.779	0.000	1.29e+05	1.81e+05
C(zipcode)[T.98074]	2.289e+05	1.03e+04	22.321	0.000	2.09e+05	2.49e+05
C(zipcode)[T.98075]	2.562e+05	1.14e+04	22.428	0.000	2.34e+05	2.79e+05
C(zipcode)[T.98077]	1.958e+05	2.94e+04	6.670	0.000	1.38e+05	2.53e+05
C(zipcode)[T.98092]	-2.162e+04	1.01e+04	-2.141	0.032	-4.14e+04	-1823.879
C(zipcode)[T.98102]	4.568e+05	1.32e+04	34.676	0.000	4.31e+05	4.83e+05
C(zipcode)[T.98103]	3.563e+05	8381.081	42.511	0.000	3.4e+05	3.73e+05
C(zipcode)[T.98105]	4.529e+05	1.04e+04	43.502	0.000	4.33e+05	4.73e+05
C(zipcode)[T.98106]	1.529e+05	9394.093	16.273	0.000	1.34e+05	1.71e+05
C(zipcode)[T.98107]	3.46e+05	9727.831	35.569	0.000	3.27e+05	3.65e+05
C(zipcode)[T.98108]	1.432e+05	1.06e+04	13.448	0.000	1.22e+05	1.64e+05
C(zipcode)[T.98109]	4.861e+05	1.29e+04	37.791	0.000	4.61e+05	5.11e+05
C(zipcode)[T.98112]	5.221e+05	1.02e+04	51.053	0.000	5.02e+05	5.42e+05
C(zipcode)[T.98115]	3.574e+05	8433.391	42.374	0.000	3.41e+05	3.74e+05
C(zipcode)[T.98116]	3.401e+05	9326.051	36.469	0.000	3.22e+05	3.58e+05
C(zipcode)[T.98117]	3.505e+05	8492.311	41.278	0.000	3.34e+05	3.67e+05
C(zipcode)[T.98118]	2.031e+05	8627.299	23.542	0.000	1.86e+05	2.2e+05
C(zipcode)[T.98119]	4.721e+05	1.09e+04	43.368	0.000	4.51e+05	4.93e+05
C(zipcode)[T.98122]	3.463e+05	9578.762	36.153	0.000	3.28e+05	3.65e+05
C(zipcode)[T.98125]	2.214e+05	9086.817	24.363	0.000	2.04e+05	2.39e+05
C(zipcode)[T.98126]	2.429e+05	9151.685	26.537	0.000	2.25e+05	2.61e+05
C(zipcode)[T.98133]	1.795e+05	8766.169	20.479	0.000	1.62e+05	1.97e+05
C(zipcode)[T.98136]	2.958e+05	9832.800	30.085	0.000	2.77e+05	3.15e+05
C(zipcode)[T.98144]	2.885e+05	9244.885	31.202	0.000	2.7e+05	3.07e+05
C(zipcode)[T.98146]	1.523e+05	1.01e+04	15.115	0.000	1.33e+05	1.72e+05
C(zipcode)[T.98148]	8.49e+04	1.71e+04	4.969	0.000	5.14e+04	1.18e+05
C(zipcode)[T.98155]	1.762e+05	9384.587	18.779	0.000	1.58e+05	1.95e+05
C(zipcode)[T.98166]	1.338e+05	1.19e+04	11.222	0.000	1.1e+05	1.57e+05
C(zipcode)[T.98168]	9.239e+04	1.07e+04	8.611	0.000	7.14e+04	1.13e+05
C(zipcode)[T.98177]	2.499e+05	1.07e+04	23.277	0.000	2.29e+05	2.71e+05
C(zipcode)[T.98178]	9.114e+04	1.02e+04	8.949	0.000	7.12e+04	1.11e+05

```
C(zipcode)[T.98188] 6.384e+04 1.29e+04 4.951 0.000 3.86e+04 8.91e+04
C(zipcode)[T.98198] 6.585e+04 1.03e+04 6.382 0.000 4.56e+04 8.61e+04
C(zipcode)[T.98199] 4.074e+05 9532.378 42.741 0.000 3.89e+05 4.26e+05
    grade_sc 4.668e+04 1267.898 36.818 0.000 4.42e+04 4.92e+04
    sqft_living_sc 7.548e+04 1357.667 55.597 0.000 7.28e+04 7.81e+04
    sqft_living15_sc 2.532e+04 1362.295 18.584 0.000 2.26e+04 2.8e+04
```

Omnibus: 4323.067 Durbin-Watson: 2.000

Prob(Omnibus): 0.000 Jarque-Bera (JB): 22074.428

Skew: 1.270 Prob(JB): 0.00

Kurtosis: 8.301 Cond. No. 108.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Validate Model using the Train-Test-Split

In [109]:

```
from sklearn.model_selection import train_test_split
X = df.drop('price', axis=1)
y = df['price']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)
```

In [110]:

```
df_train = pd.DataFrame(X_train, columns = X.columns)
df_train['price'] = y_train
df_test = pd.DataFrame(X_test, columns = X.columns)
df_test['price'] = y_test
```

In [111]:

```
from sklearn.metrics import r2_score
y_hat_test = model.predict(df_test)
r2_score(y_test, y_hat_test)
```

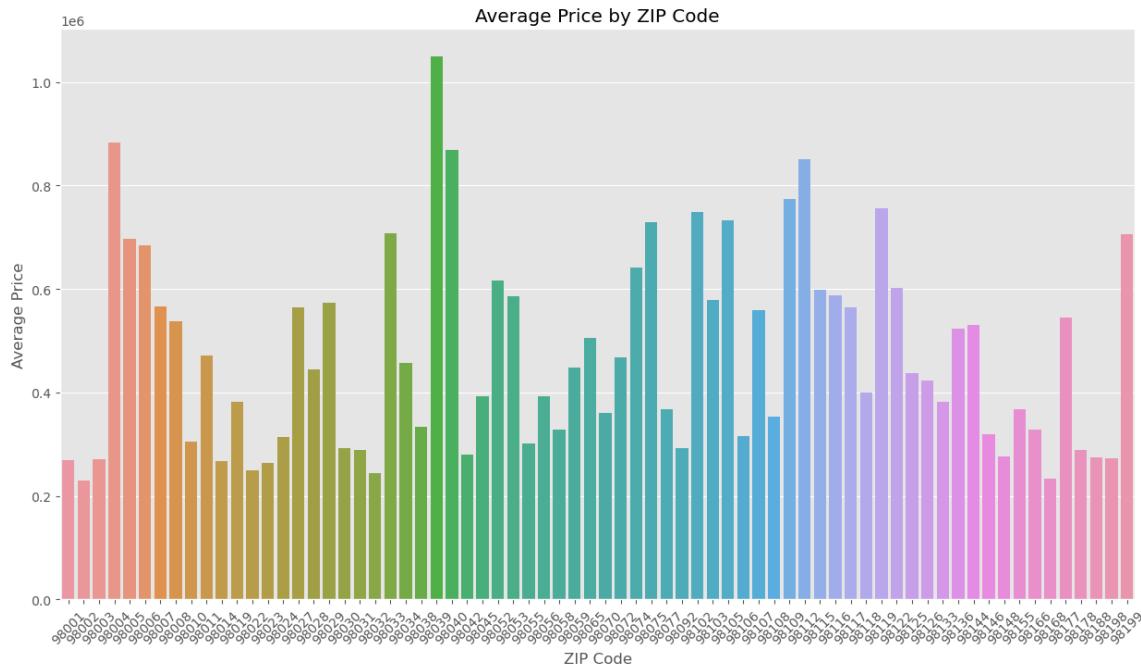
Out[111]:

0.768252174841723

Visual data to figure out the average price per zip codes

In [112]:

```
plt.figure(figsize=(15, 8))
sns.barplot(x='zipcode', y='price', data=df, errorbar=None)
plt.xticks(rotation=45)
plt.xlabel('ZIP Code')
plt.ylabel('Average Price')
plt.title('Average Price by ZIP Code')
plt.show()
```

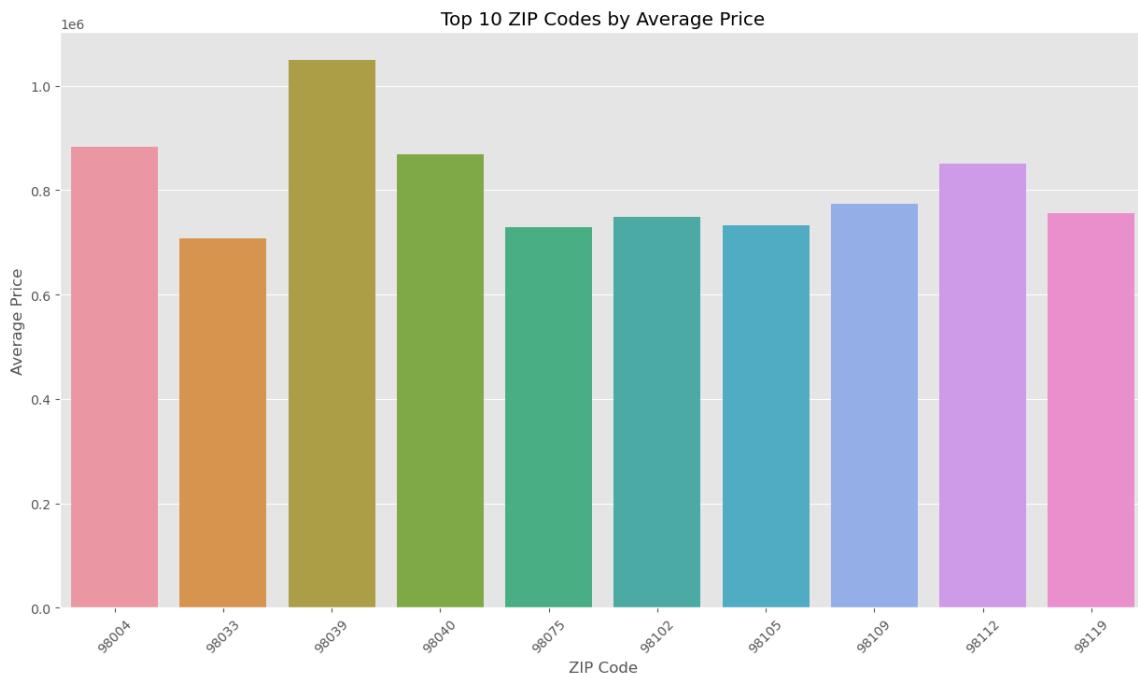


In [113]:

```
# Calculate the average price for each ZIP code
avg_price_per_zip = df.groupby('zipcode')['price'].mean().reset_index()

# Sort ZIP codes by average price in descending order and take top 10
top_10_zip = avg_price_per_zip.sort_values('price', ascending=False).head(10)

# Plotting
plt.figure(figsize=(15, 8))
sns.barplot(x='zipcode', y='price', data=top_10_zip, errorbar=None)
plt.xticks(rotation=45)
plt.xlabel('ZIP Code')
plt.ylabel('Average Price')
plt.title('Top 10 ZIP Codes by Average Price')
plt.show()
```



Analysis Summary:

Upon conducting multiple regression analyses, the fifth model demonstrated a robust R-squared value of 0.775, signifying that approximately 78% of the variability in home prices within King County can be explained by the predictors in the model. Notably, the postal code emerged as one of the most significant determinants influencing property valuation. Additional variables that were positively correlated with property prices include:

- Living Space Square Footage
- Building Grade, as per the King County Grading System
- Average Interior Square Footage of the Nearest 15 Residences

Strategic Recommendations for a Pro Home Builder Company:

To optimize the profitability and marketability of residential properties in King County, the following strategies are recommended:

- Strategic Location Selection: Prioritize property acquisition in high-value zip codes to benefit from the area's intrinsically higher property valuations.
- Optimize Living Space: Expand the square footage of the living area to align with the positive correlation observed between living space and property value.
- Quality of Construction: Employ premium-quality building materials and construction techniques to secure the highest possible building grade according to the King County Grading System. This will not only enhance the property's intrinsic value but also its appeal to potential buyers.
- Neighborhood Composition: Opt for land parcels where the surrounding 15 residences have a higher average interior square footage. This is indicative of a neighborhood's overall property value and can positively influence the value of the new development.

By implementing these targeted strategies, housing development firms can substantially increase